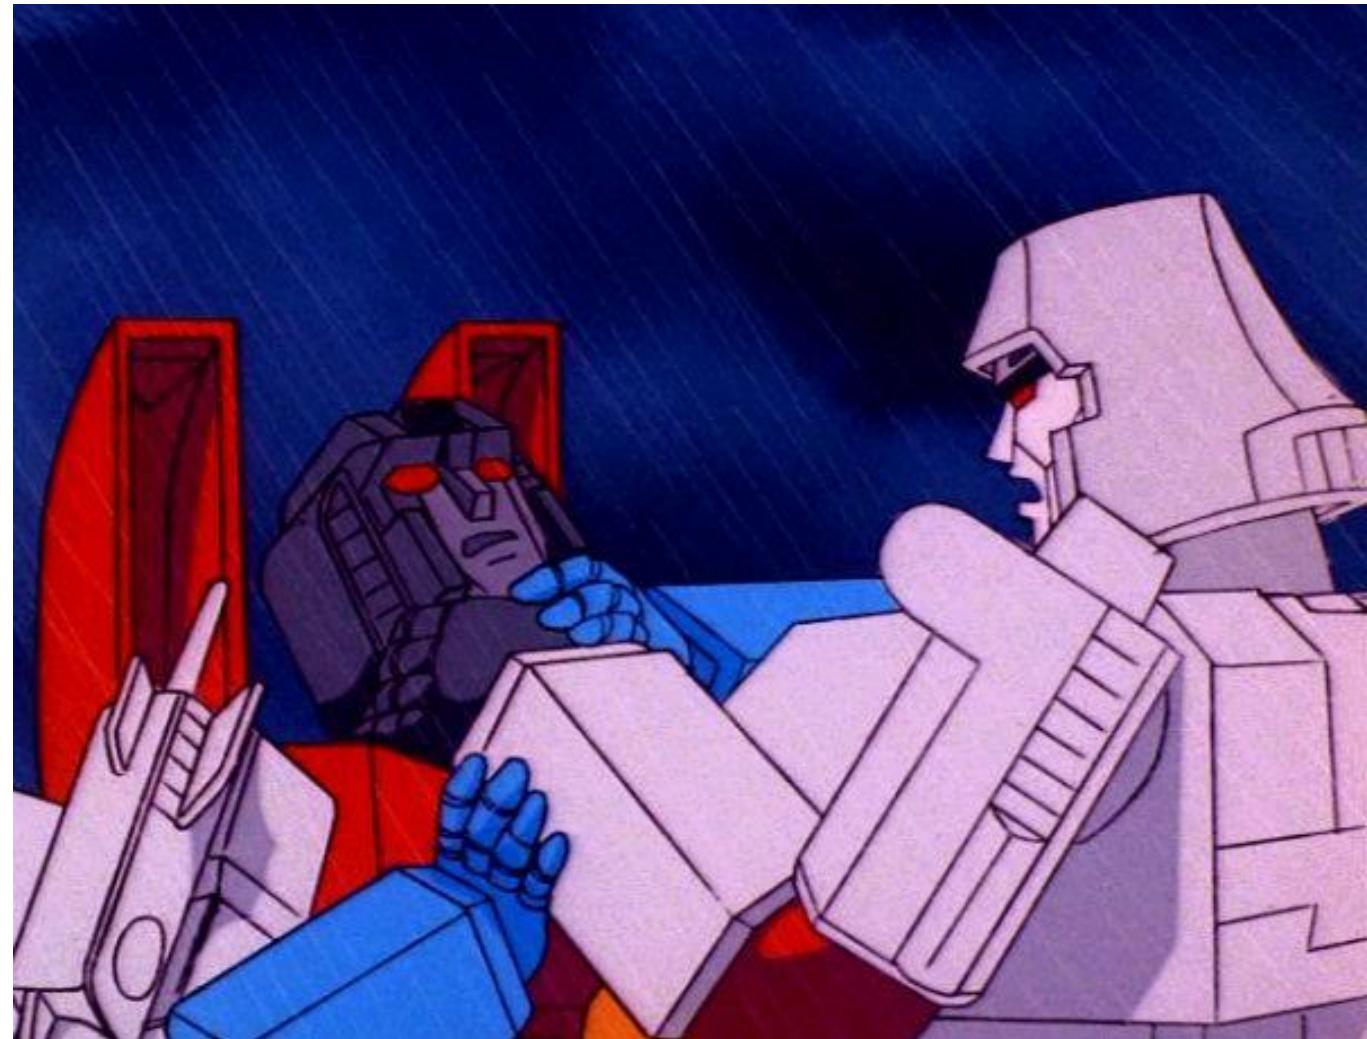


Transformers

Lecture 6





Language Modelling Task

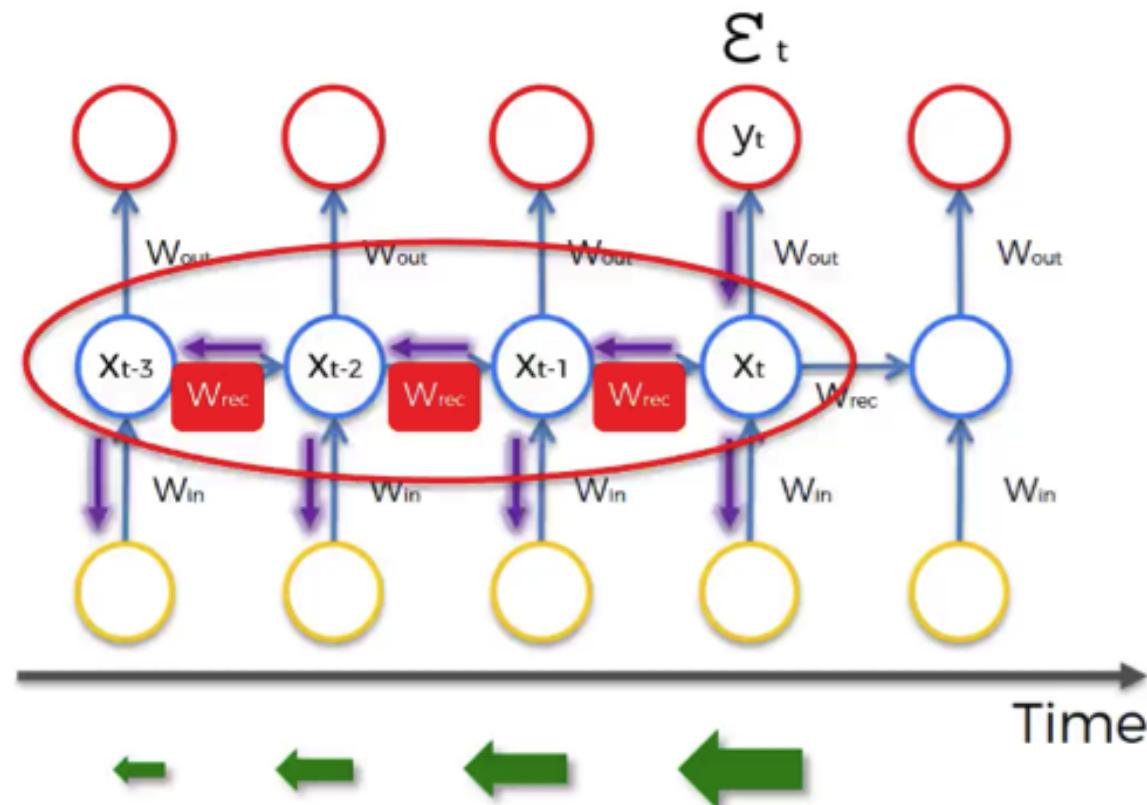
- Final goal: predict/estimate the probability of a sequence

Probability(*Some sentence over here.*)

- Actual task:
 - Predict the next word
 - MLM
- In a perfect world:
 - The RNN hidden states should be able capture all contextual information
 - Right?



The Vanishing Gradient Problem



$$\frac{\partial \mathcal{E}}{\partial \theta} = \sum_{1 \leq t \leq T} \frac{\partial \mathcal{E}_t}{\partial \theta} \quad (3)$$

$$\frac{\partial \mathcal{E}_t}{\partial \theta} = \sum_{1 \leq k \leq t} \left(\frac{\partial \mathcal{E}_t}{\partial \mathbf{x}_t} \frac{\partial \mathbf{x}_t}{\partial \mathbf{x}_k} \frac{\partial^+ \mathbf{x}_k}{\partial \theta} \right) \quad (4)$$

$$\frac{\partial \mathbf{x}_t}{\partial \mathbf{x}_k} = \prod_{t \geq i > k} \frac{\partial \mathbf{x}_i}{\partial \mathbf{x}_{i-1}} = \prod_{t \geq i > k} \mathbf{W}_{rec}^T diag(\sigma'(\mathbf{x}_{i-1})) \quad (5)$$

$W_{rec} \sim \text{small} \rightarrow$ Vanishing
 $W_{rec} \sim \text{large} \rightarrow$ Exploding

Formula Source: Razvan Pascanu et al. (2013)

Again, we need more! What of larger semantic units?

- How can we know when larger units are similar in meaning?
 - *CTV News*: Poilievre-led attempt to bring down Trudeau minority over carbon tax fails.
 - *CBC News*: Liberals survive non-confidence vote on carbon tax with Bloc, NDP backing.
 - *The Beaverton*: Co-worker that everyone hates surprised he can't get colleagues to do what he wants.



NATIONAL - 2 WEEKS AGO

Co-worker that everyone hates surprised he can't get colleagues to do what he wants

OTTAWA – Local man Pierre Poilievre, an employee at an Ottawa small business named the House of Commons, was surprised that none of the colleagues who despise him were willing to support hi...



SHARE

RNN & next word prediction: Not good compositional representation

- Next word prediction:

$$P(t_i | t_1, t_2, \dots, t_{i-1})$$

- The hidden state i is encoding information of everything from the beginning (index 0) to the very end (index i).
- We want some bigger semantic units
 - Poilievre-led attempt to **bring down Trudeau minority over carbon tax** fails.
- Some hacks may work, but not really

Again, we need more! What of larger semantic units?

- How can we know when larger units are similar in meaning?
 - *CTV News*: Poilievre-led attempt to bring down Trudeau minority over carbon tax fails.
 - *CBC News*: Liberals survive non-confidence vote on carbon tax with Bloc, NDP backing.
 - *The Beaverton*: Co-worker that everyone hates surprised he can't get colleagues to do what he wants.

People interpret the meaning of larger text units – entities, descriptive terms, facts, arguments, stories – by **semantic composition** of smaller elements.

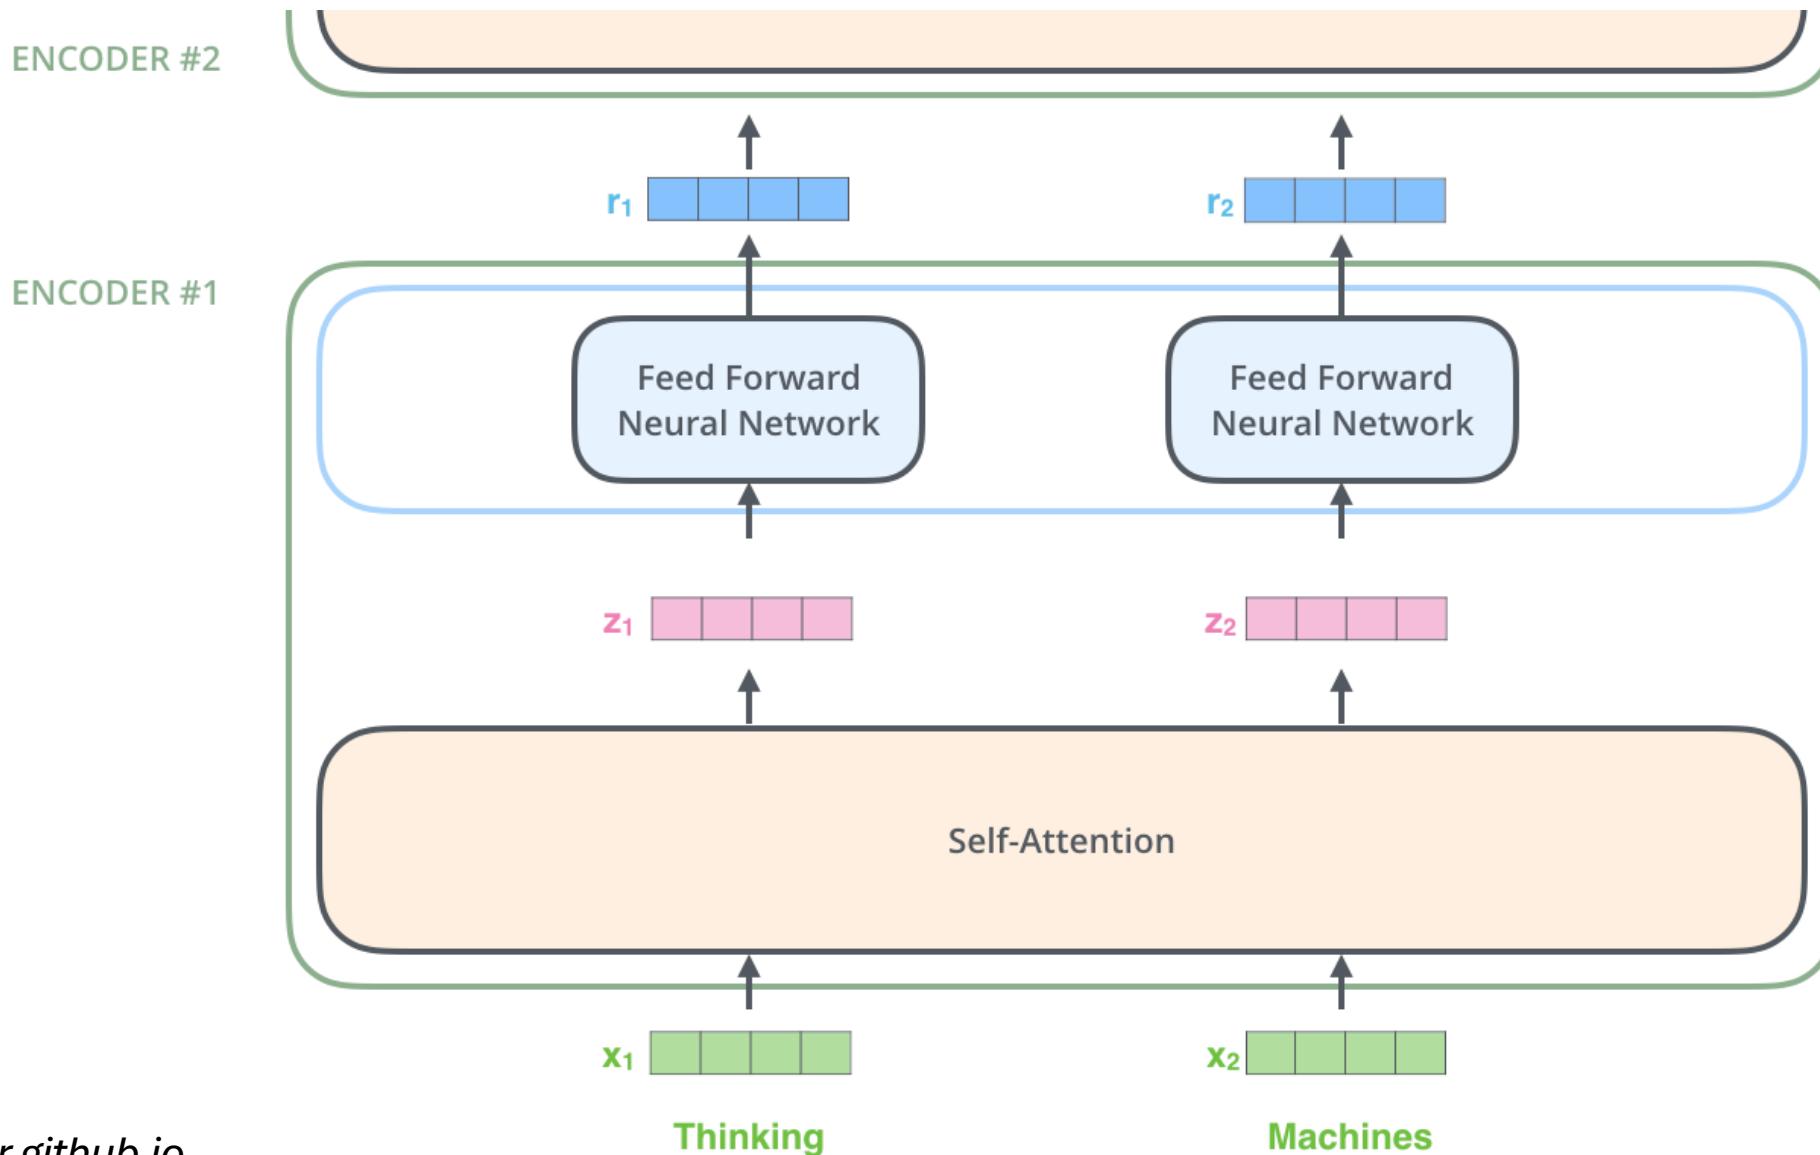
Beyond Word2Vec and RNN-LMs

- Mitigation #1:
 - Long short-term memory (LSTM).
 - Won't cover in this class, but covered DL4NLP.
- Mitigation #2:
 - Attention Mechanism -> Transformer

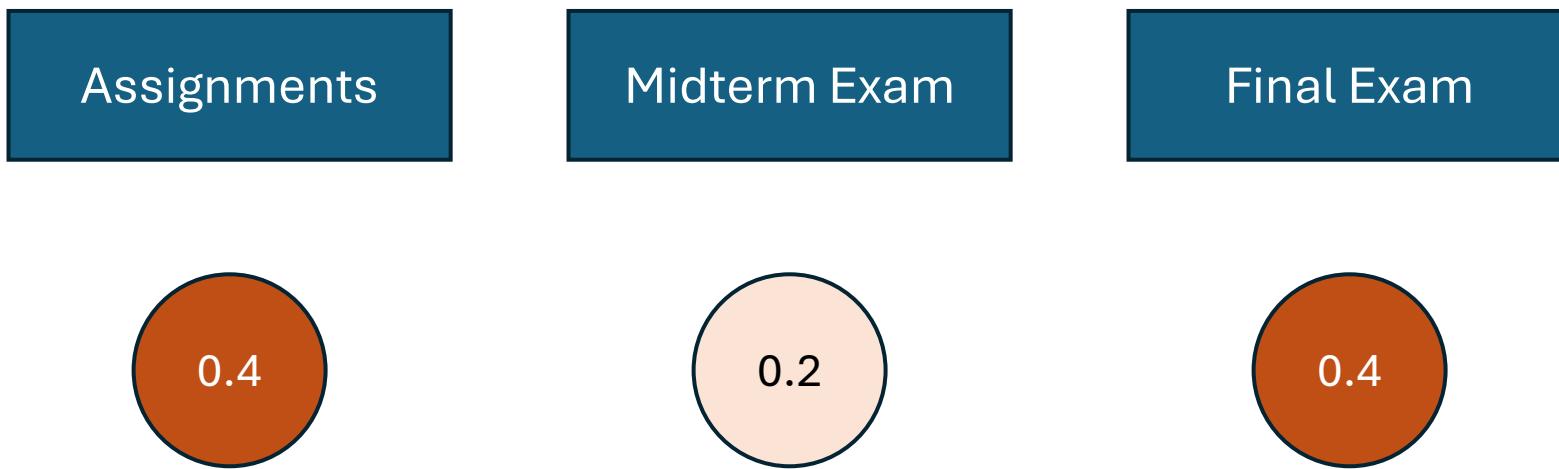
LSTM

- (Won't appear on your exam)
- Basic idea:
 - Reset the hidden state once a while.
 - When to know how to reset? Train a separate model (actually, a part of the model) to do it.

Transformers

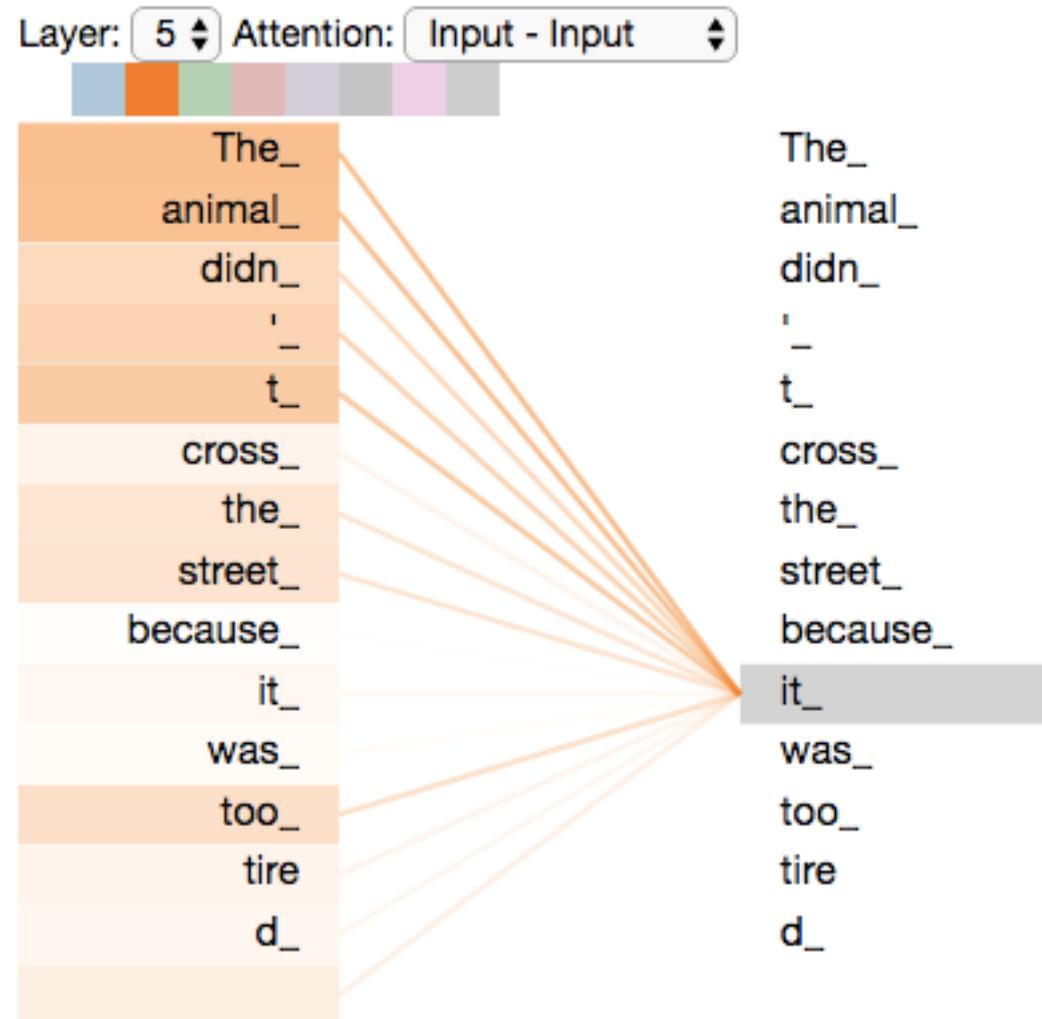


Attention Mechanism

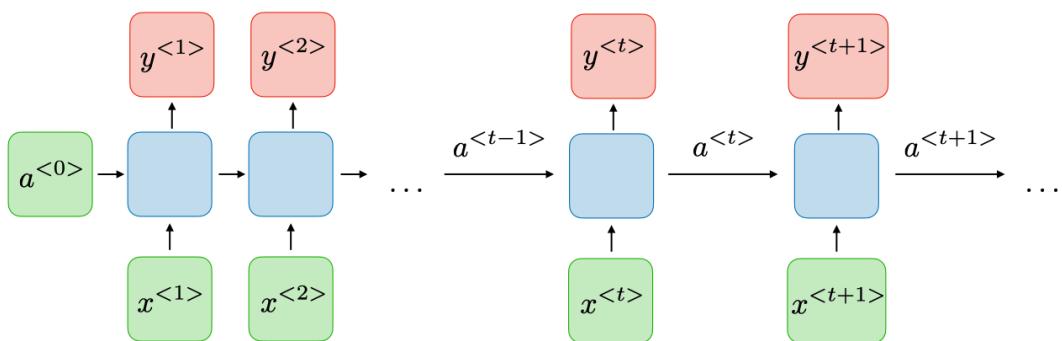


$$\text{Total} = \boxed{0.75} * A + \boxed{0.1} * Q + \boxed{0.15} * E$$

Attention Mechanism

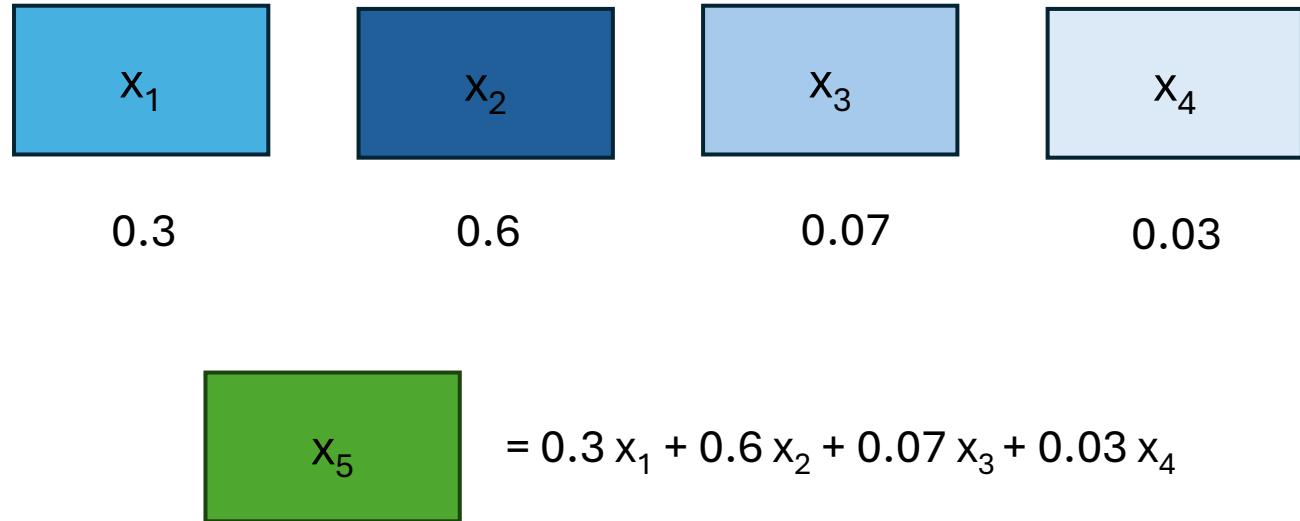


RNN vs. Attention



Recurrent neural network.
Almost “recursive.”

$$h_i = f(h_{i-1} + x_i)$$



$$x_i = \sum_j a_j x_j$$

????

$$h_i = f(x_i)$$

Recall GloVe

Encoding meaning in vectors

- How can we capture ratios of co-occurrence probability components in a word vector space?
- Solution:

- Log-bilinear model:

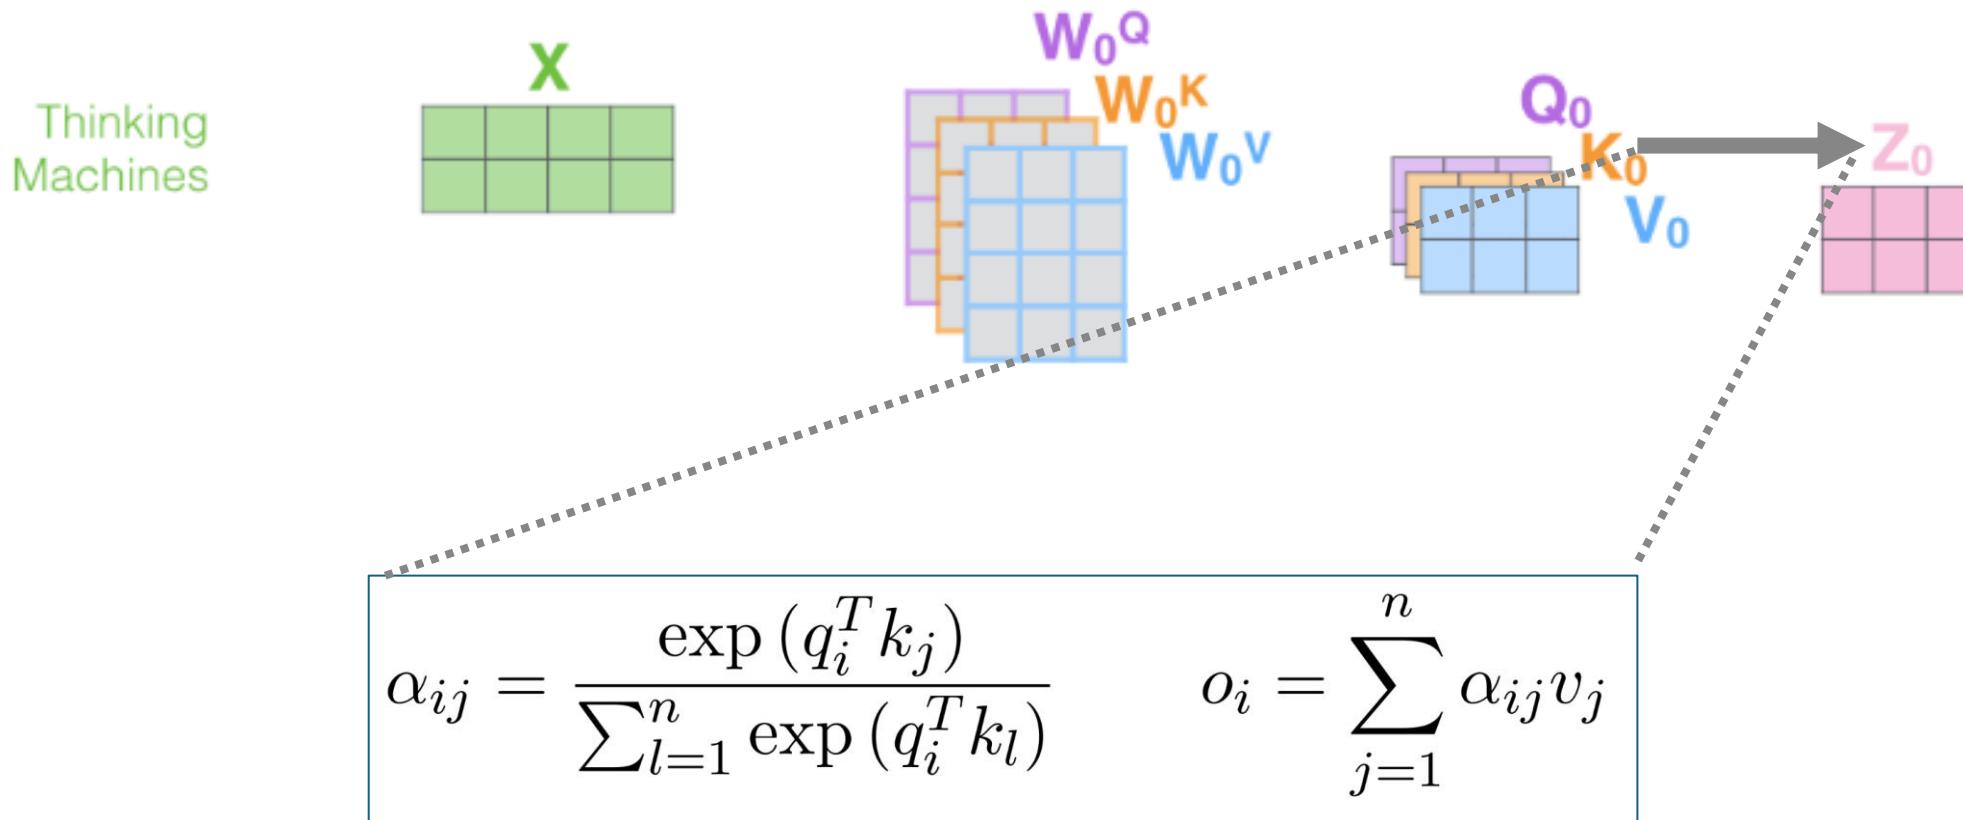
$$w_i \cdot w_j = \log P(i|j)$$

- with vector differences:

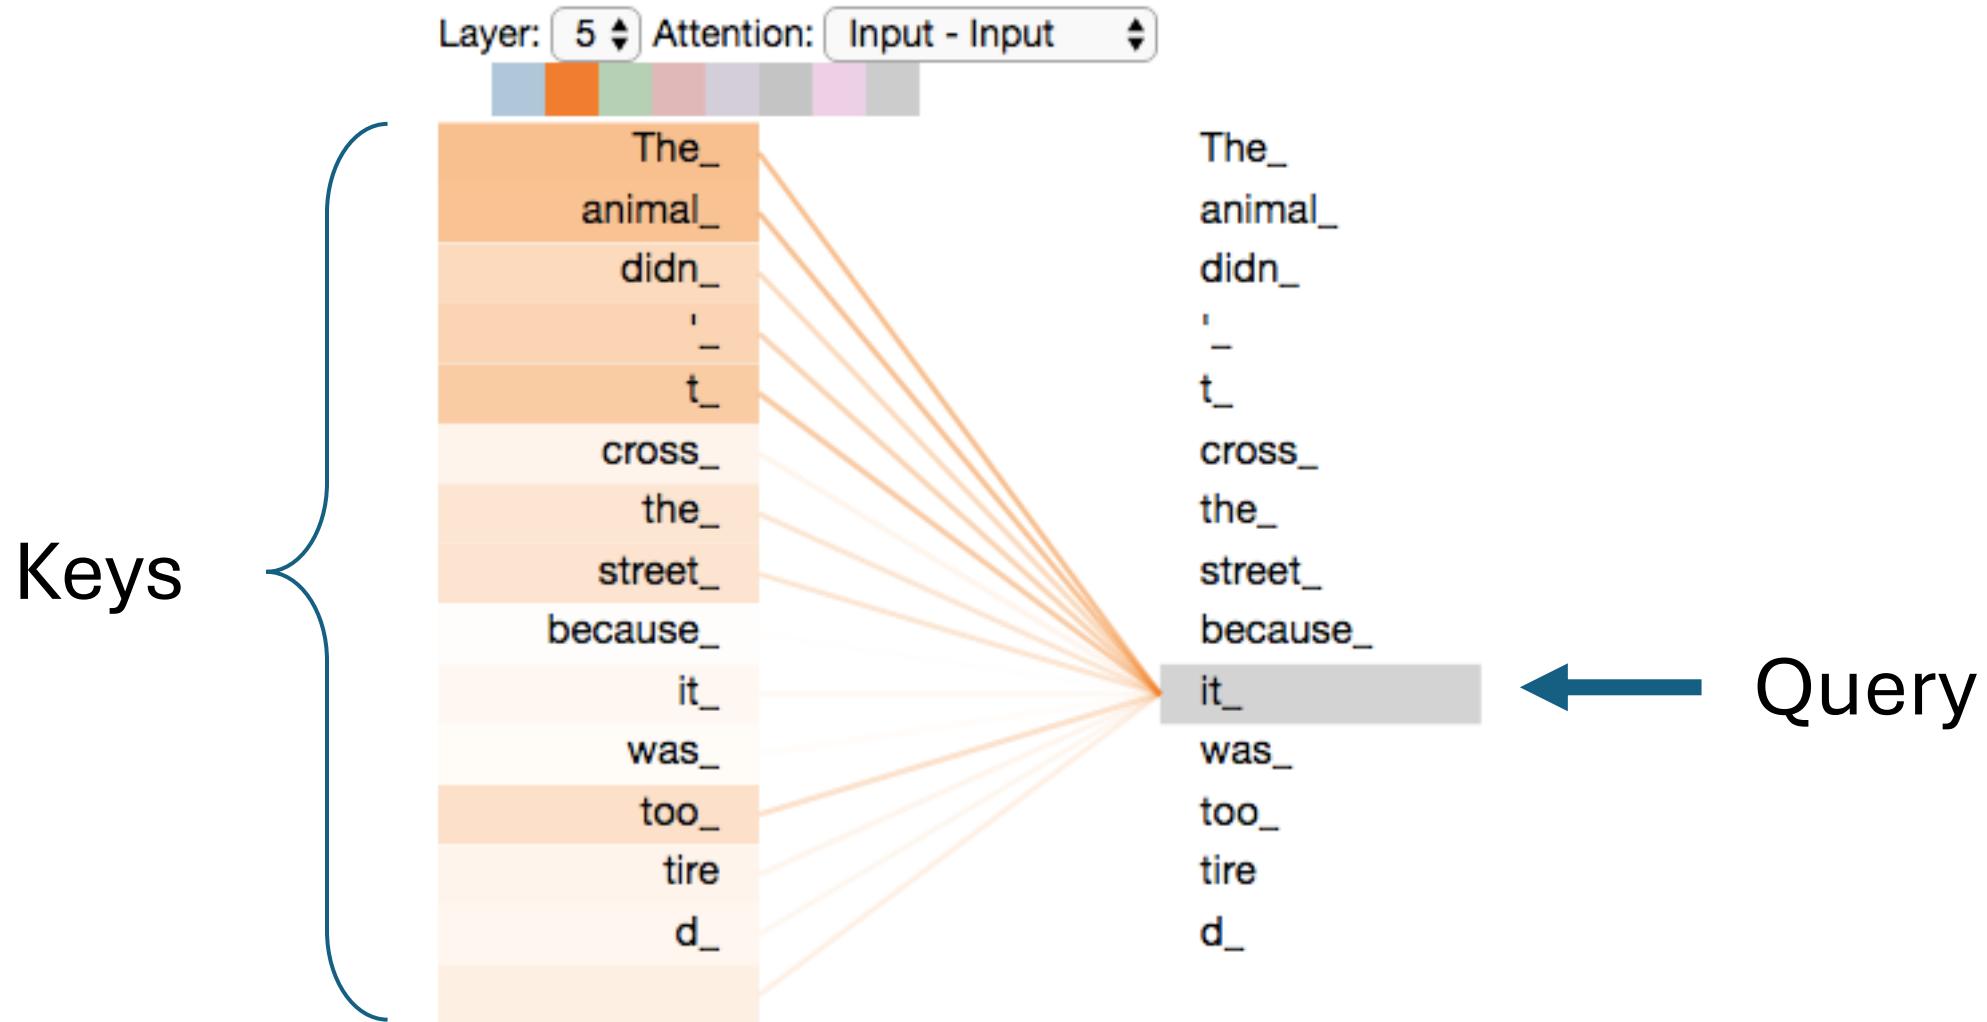
$$w_x \cdot (w_a - w_b) = \log \frac{P(x|a)}{P(x|b)}$$

Pennington et al. (2014)

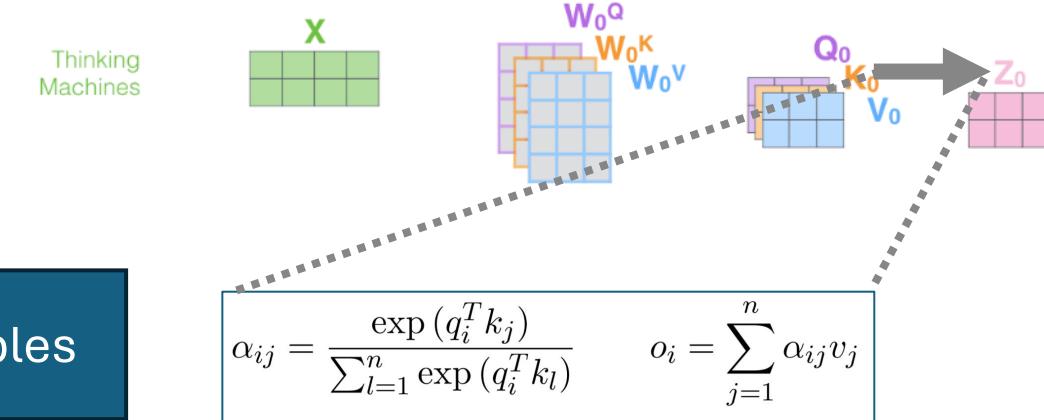
Self Attention



Attention Mechanism



Self Attention



I like eating apples

What is $o(\text{like})$? i.e., the attention score from like to the sentence.

query

Q_2
“like”

key

K1 “I”	K2 “like”	K3 “eating”	K4 “apples”
-----------	--------------	----------------	----------------

attention
weights*

14.2	18.1	10.3	7.9
------	------	------	-----

attention
scores* (o)

0.3	0.6	0.07	0.03
-----	-----	------	------

*: made-up numbers, not real.

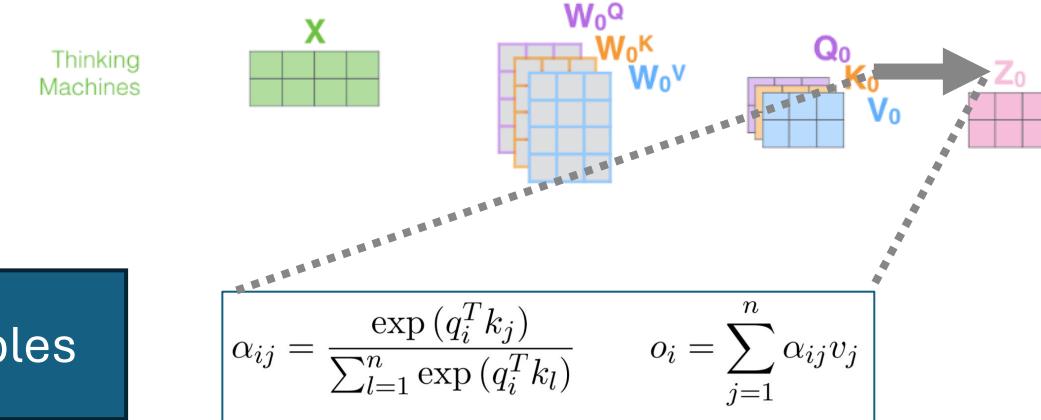
$$X \times W^Q = Q$$

$$X \times W^K = K$$

$$Q \times K^T$$

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}} \right)$$

Self Attention



I like eating apples

What is $z(\text{like})$?

attention scores* (o) 0.3 0.6 0.07 0.03

values

V1 “I”	V2 “like”	V3 “eating”	V4 “apples”
-----------	--------------	----------------	----------------

Weighted sum

Z2
“like”

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}} \right)$$
$$X \times W^V = V$$

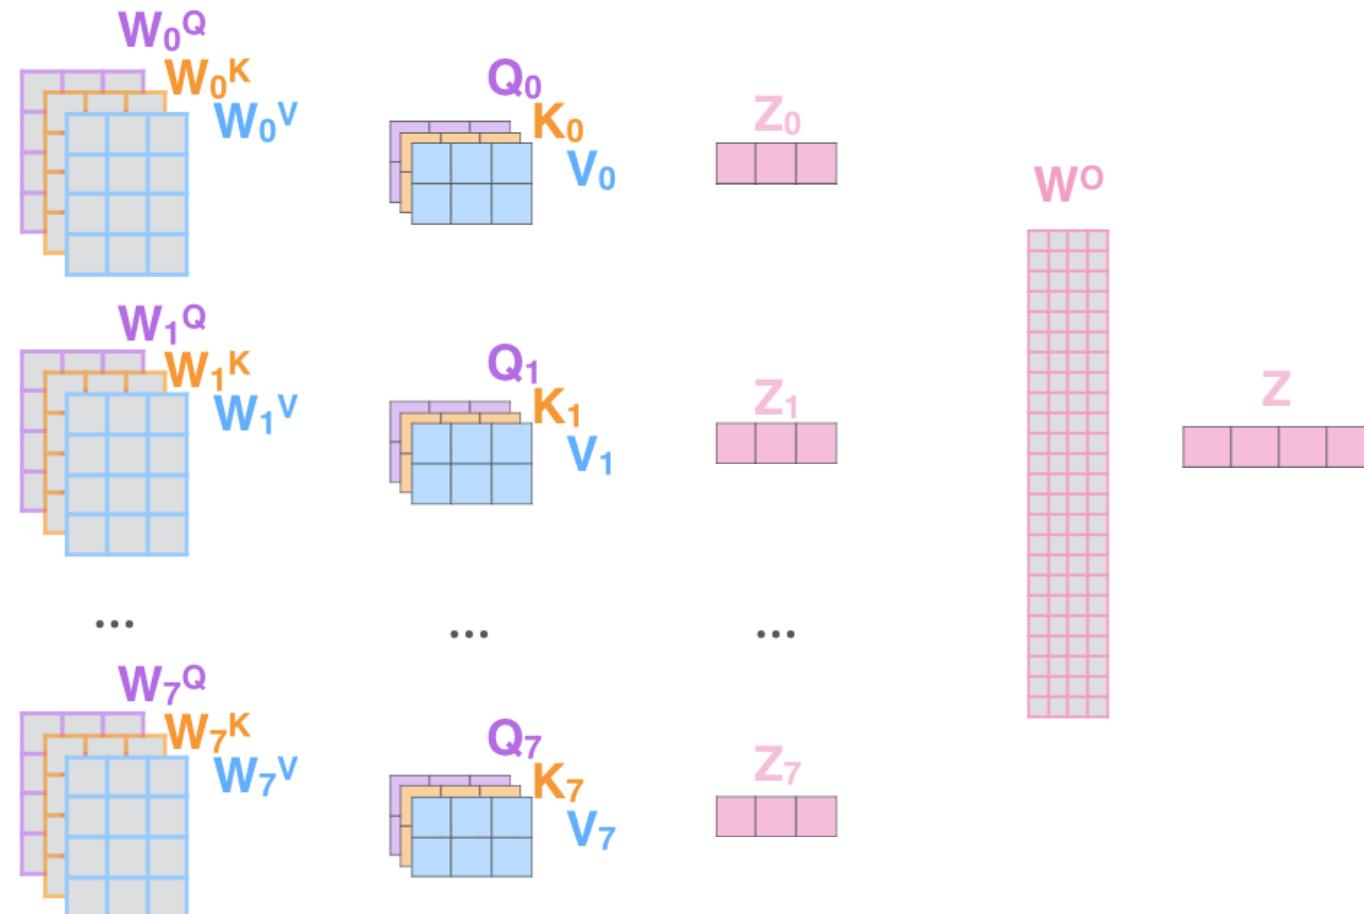
*: made-up numbers, not real.

Multi-Head Self Attention

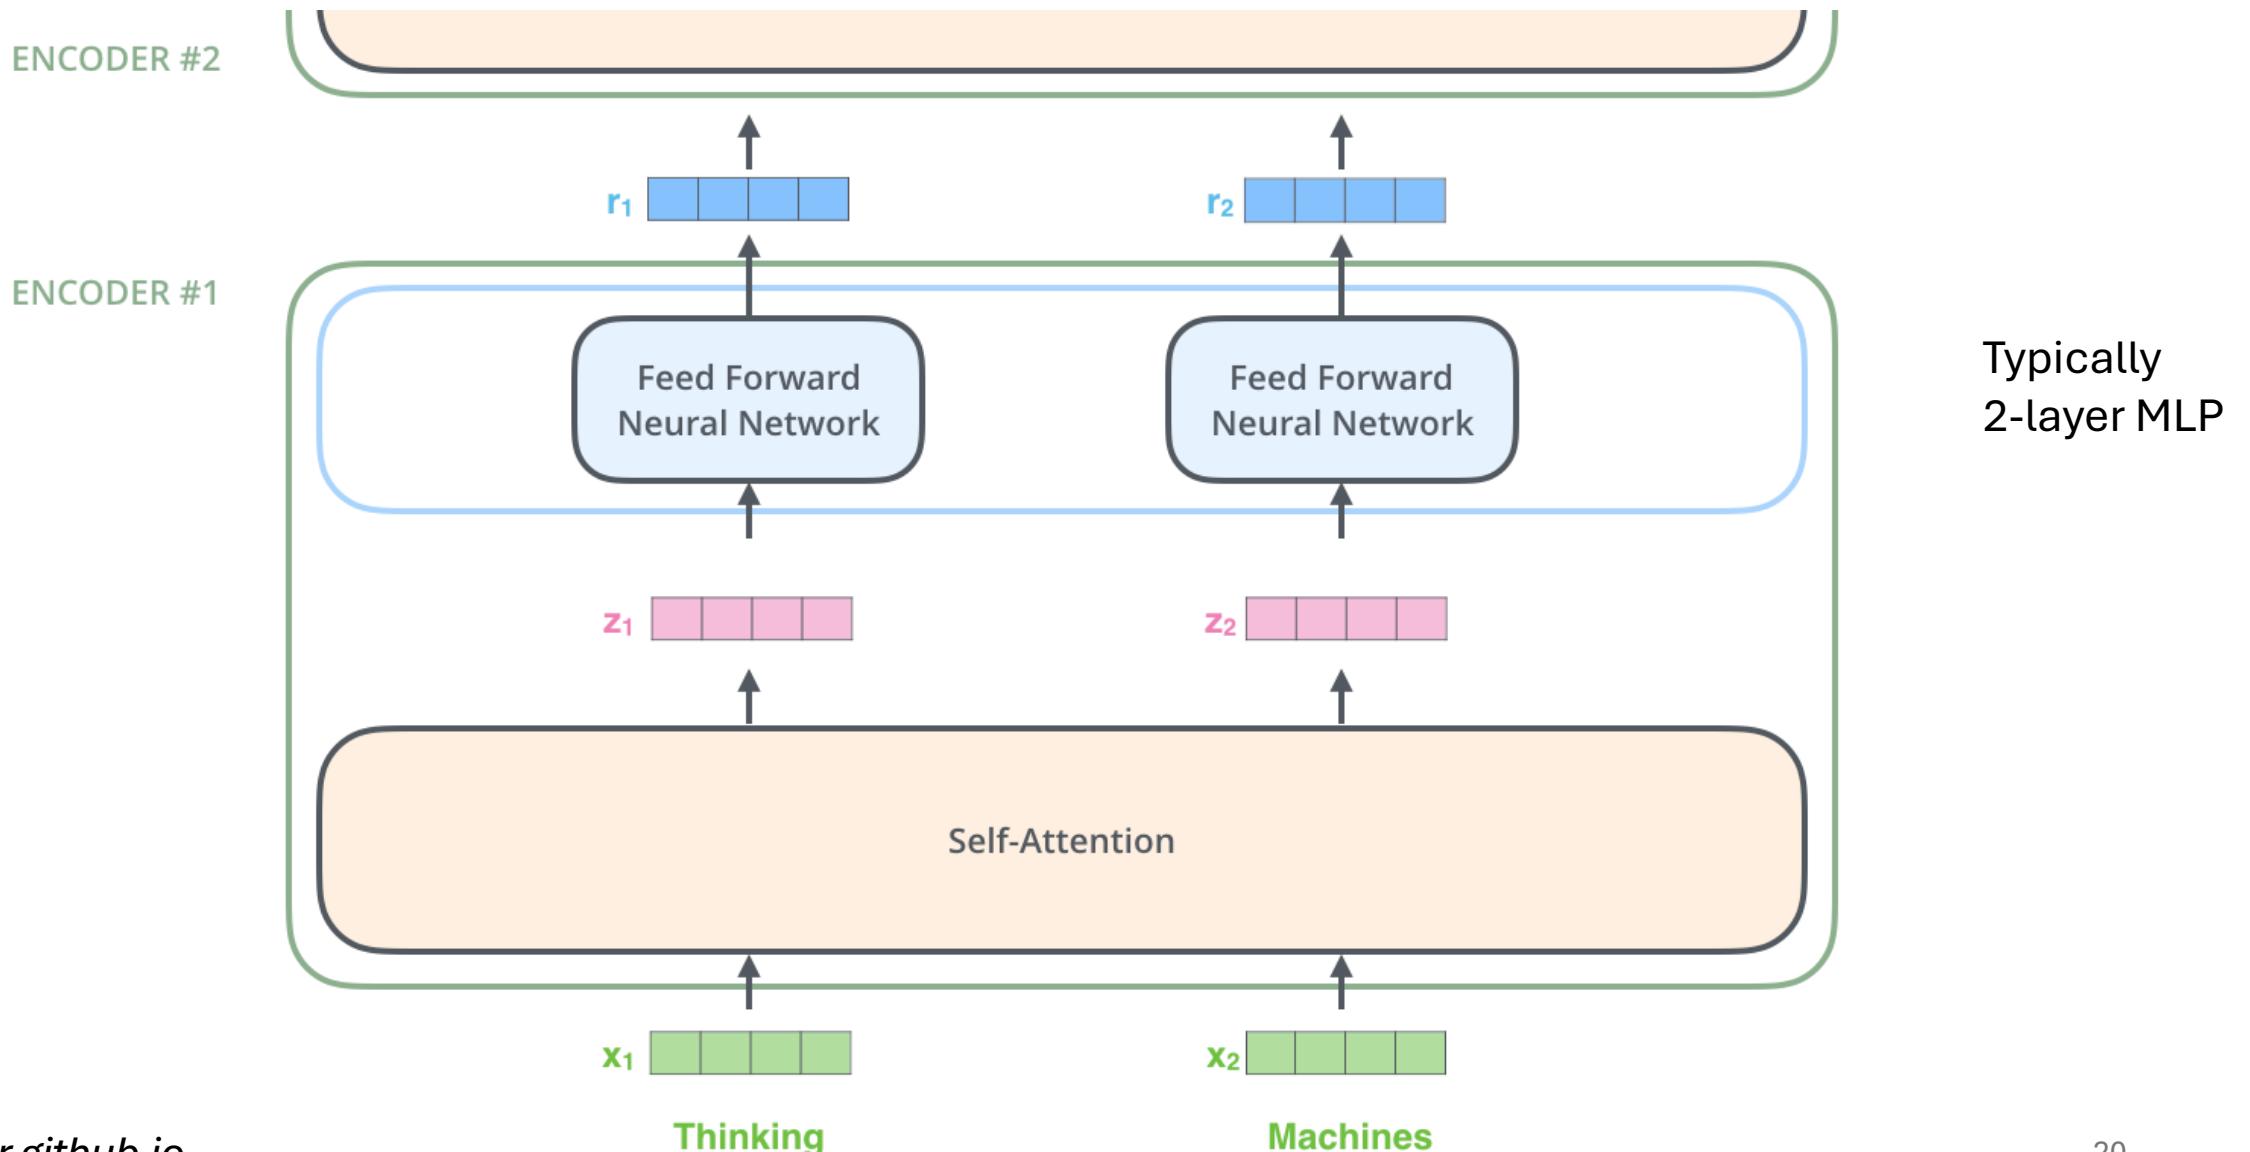
- 1) This is our input sentence* each word*
- 2) We embed each word*
- 3) Split into 8 heads. We multiply X or R with weight matrices
- 4) Calculate attention using the resulting $Q/K/V$ matrices
- 5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer



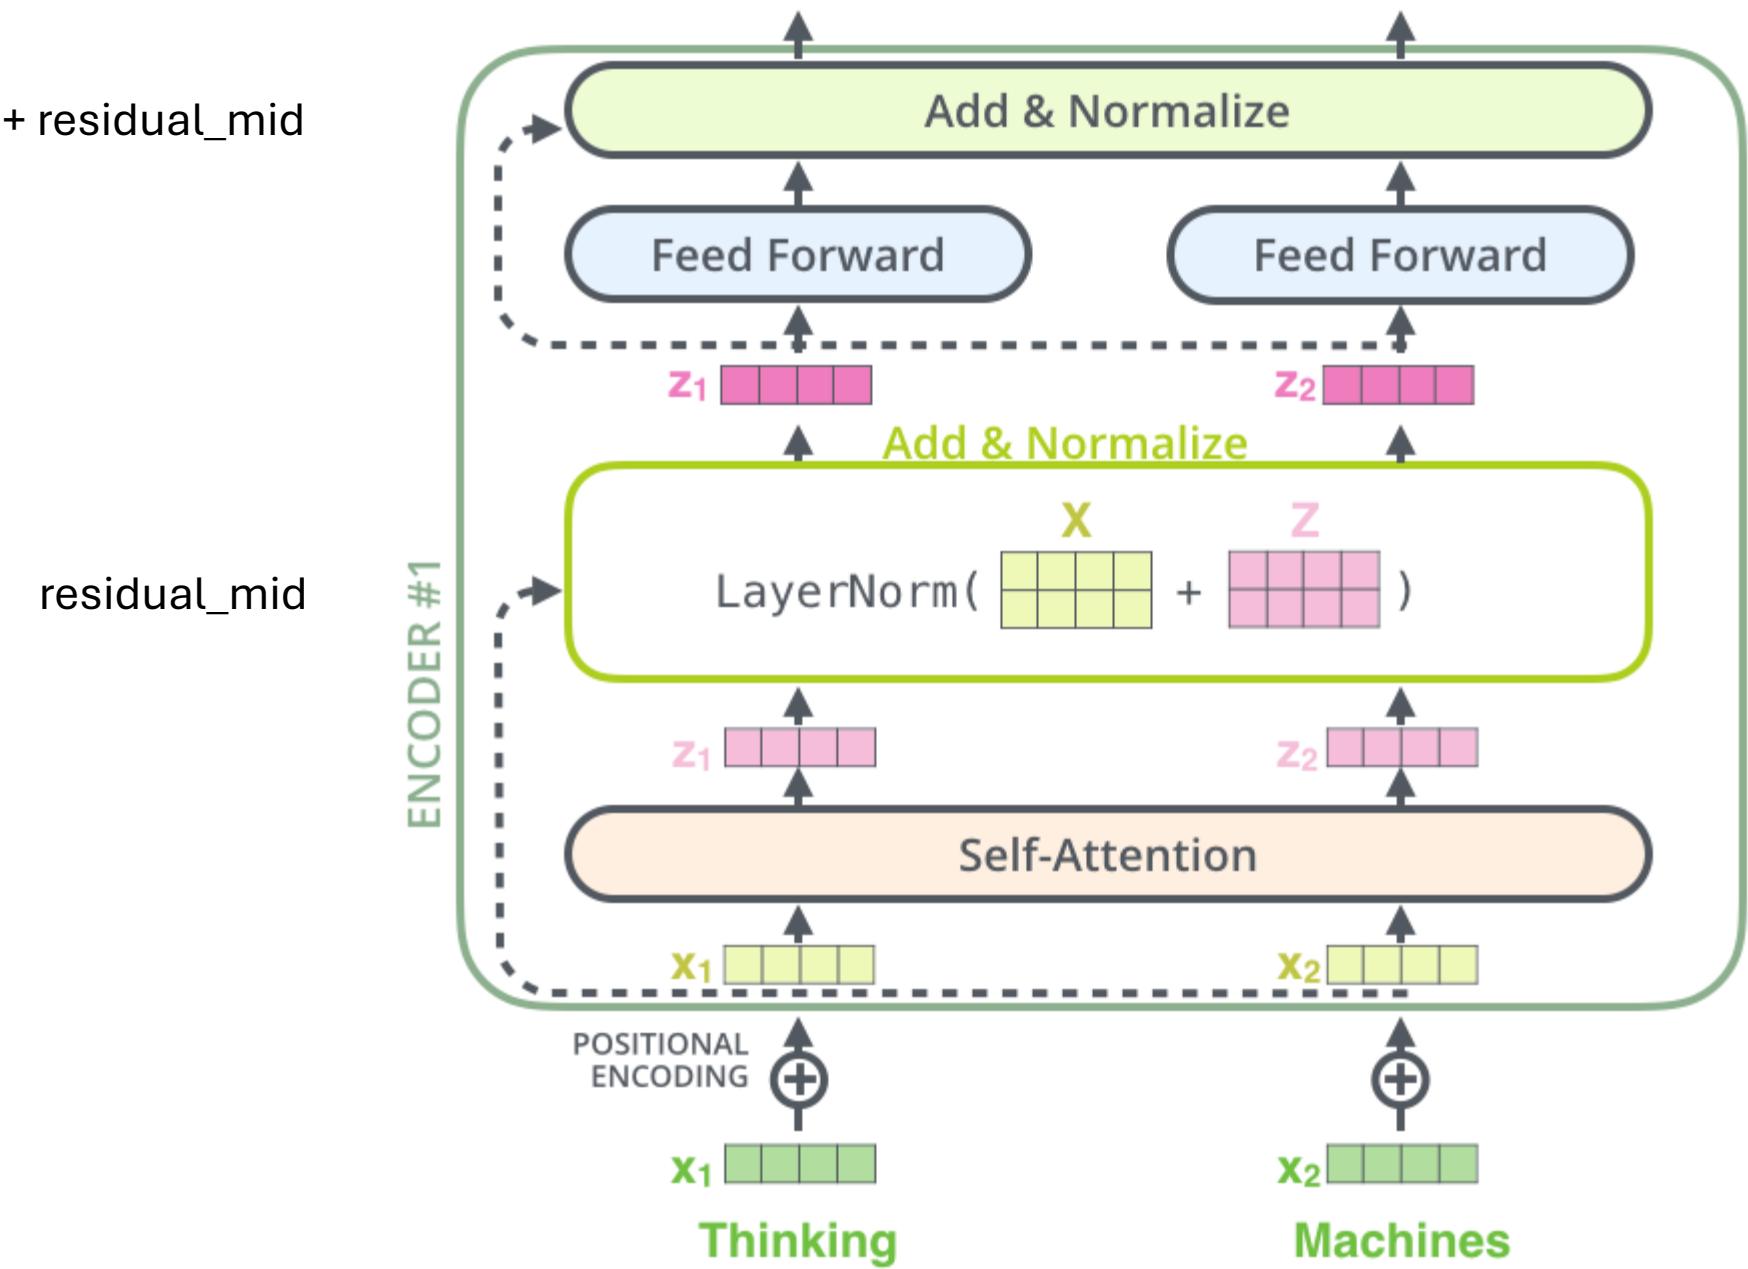
* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



Transformers



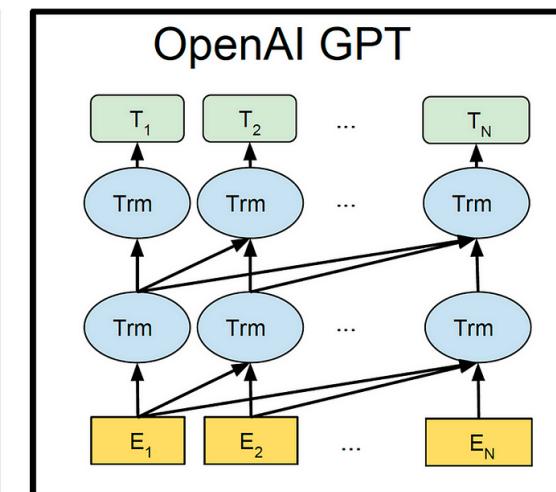
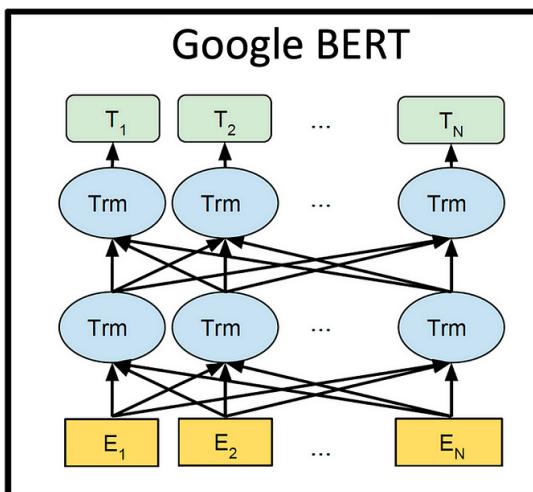
Residuals

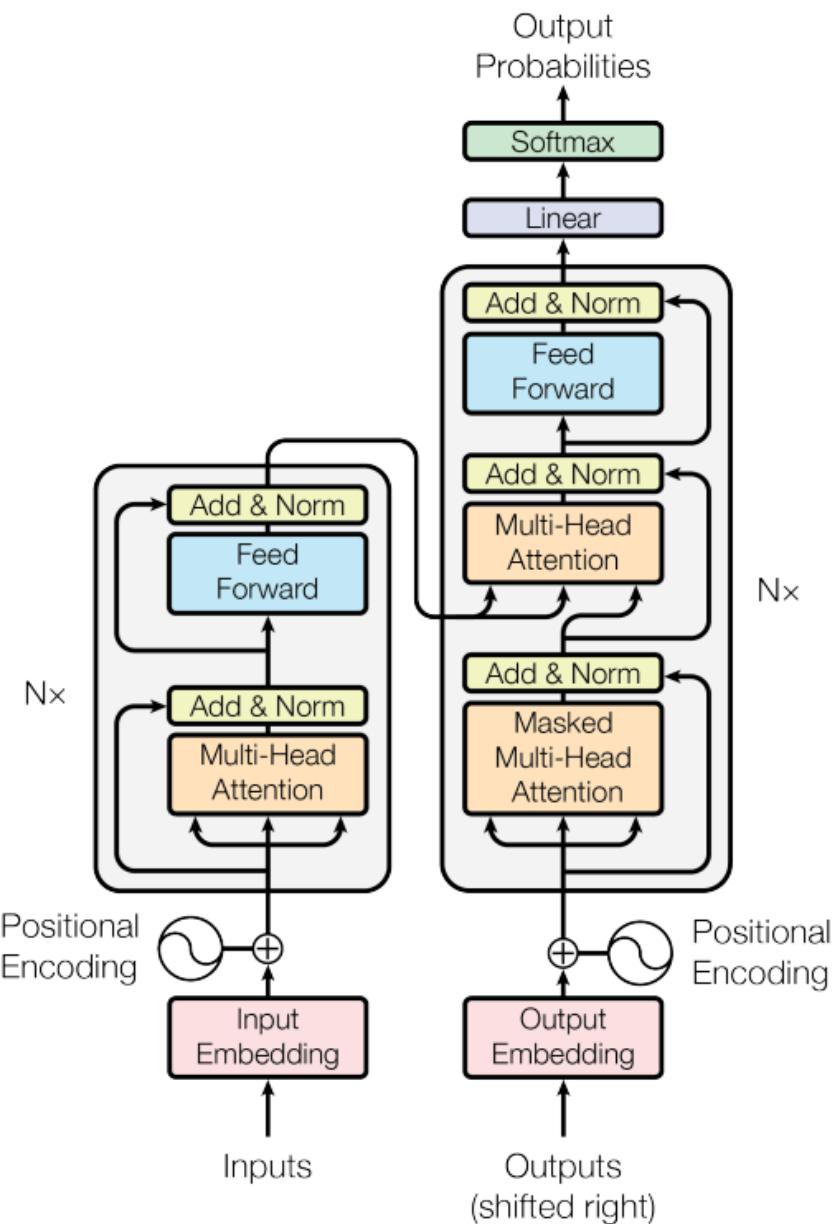


“Ogres are like onions

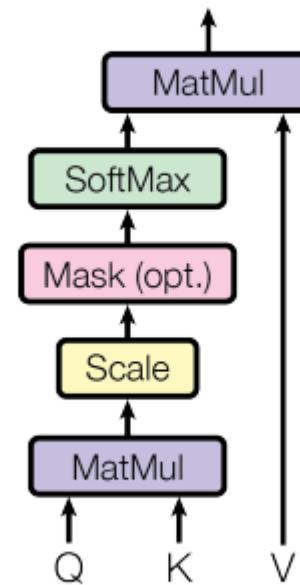
they have layers”

- Shrek

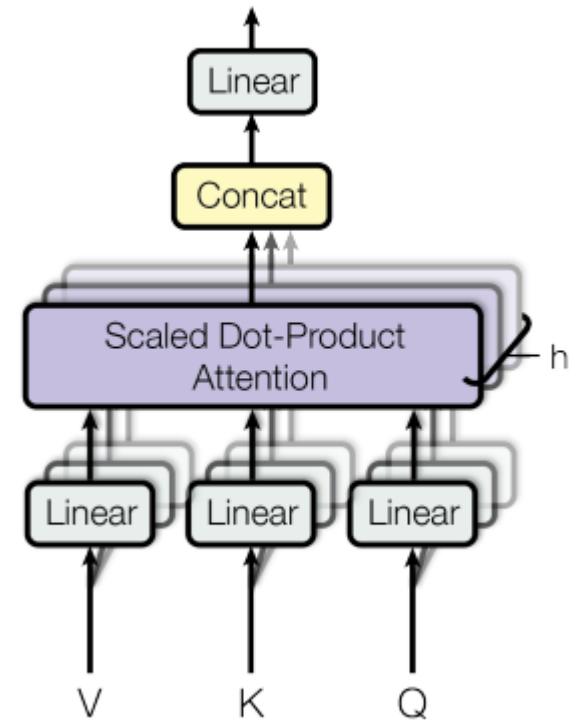




Scaled Dot-Product Attention



Multi-Head Attention



Position Encoding

Input	[CLS]	my	dog	is	cute	[SEP]	he	likes	play	#ing	[SEP]
Token Embeddings	E _[CLS]	E _{my}	E _{dog}	E _{is}	E _{cute}	E _[SEP]	E _{he}	E _{likes}	E _{play}	E _{#ing}	E _[SEP]
Segment Embeddings	E _A	E _A	E _A	E _A	E _A	E _A	E _B	E _B	E _B	E _B	E _B
Position Embeddings	E ₀	E ₁	E ₂	E ₃	E ₄	E ₅	E ₆	E ₇	E ₈	E ₉	E ₁₀

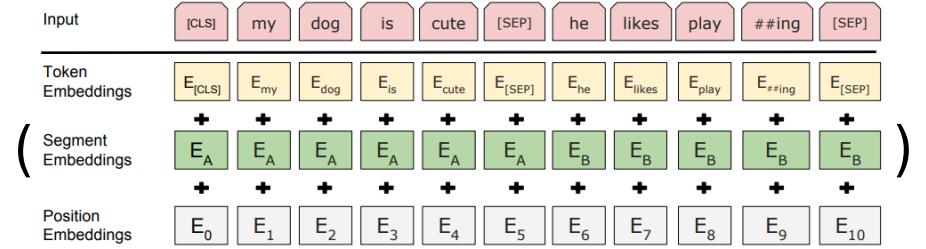
$$\vec{p}_t = \begin{bmatrix} \sin(\omega_1 \cdot t) \\ \cos(\omega_1 \cdot t) \\ \sin(\omega_2 \cdot t) \\ \cos(\omega_2 \cdot t) \\ \vdots \\ \sin(\omega_{d/2} \cdot t) \\ \cos(\omega_{d/2} \cdot t) \end{bmatrix}_{d \times 1}$$

where

$$\vec{p}_t^{(i)} = f(t)^{(i)} := \begin{cases} \sin(\omega_k \cdot t), & \text{if } i = 2k \\ \cos(\omega_k \cdot t), & \text{if } i = 2k + 1 \end{cases}$$

$$\omega_k = \frac{1}{10000^{2k/d}}$$

Position Encoding



- Encodings of any two distinct positions are distinct
- Each position maps to only one encoding
- Test sentences may be longer than training
- For any fixed offset k , $\text{PE}(t+k)$ can be represented as a linear function of $\text{PE}(t)$.

$$\vec{p}_t = \begin{bmatrix} \sin(\omega_1 \cdot t) \\ \cos(\omega_1 \cdot t) \\ \sin(\omega_2 \cdot t) \\ \cos(\omega_2 \cdot t) \\ \vdots \\ \sin(\omega_{d/2} \cdot t) \\ \cos(\omega_{d/2} \cdot t) \end{bmatrix}_{d \times 1}$$

where

$$\vec{p}_t^{(i)} = f(t)^{(i)} := \begin{cases} \sin(\omega_k \cdot t), & \text{if } i = 2k \\ \cos(\omega_k \cdot t), & \text{if } i = 2k + 1 \end{cases}$$

$$\omega_k = \frac{1}{10000^{2k/d}}$$

Sinusoidal: Linear Relative Offset

On linear relation between two encoded positions

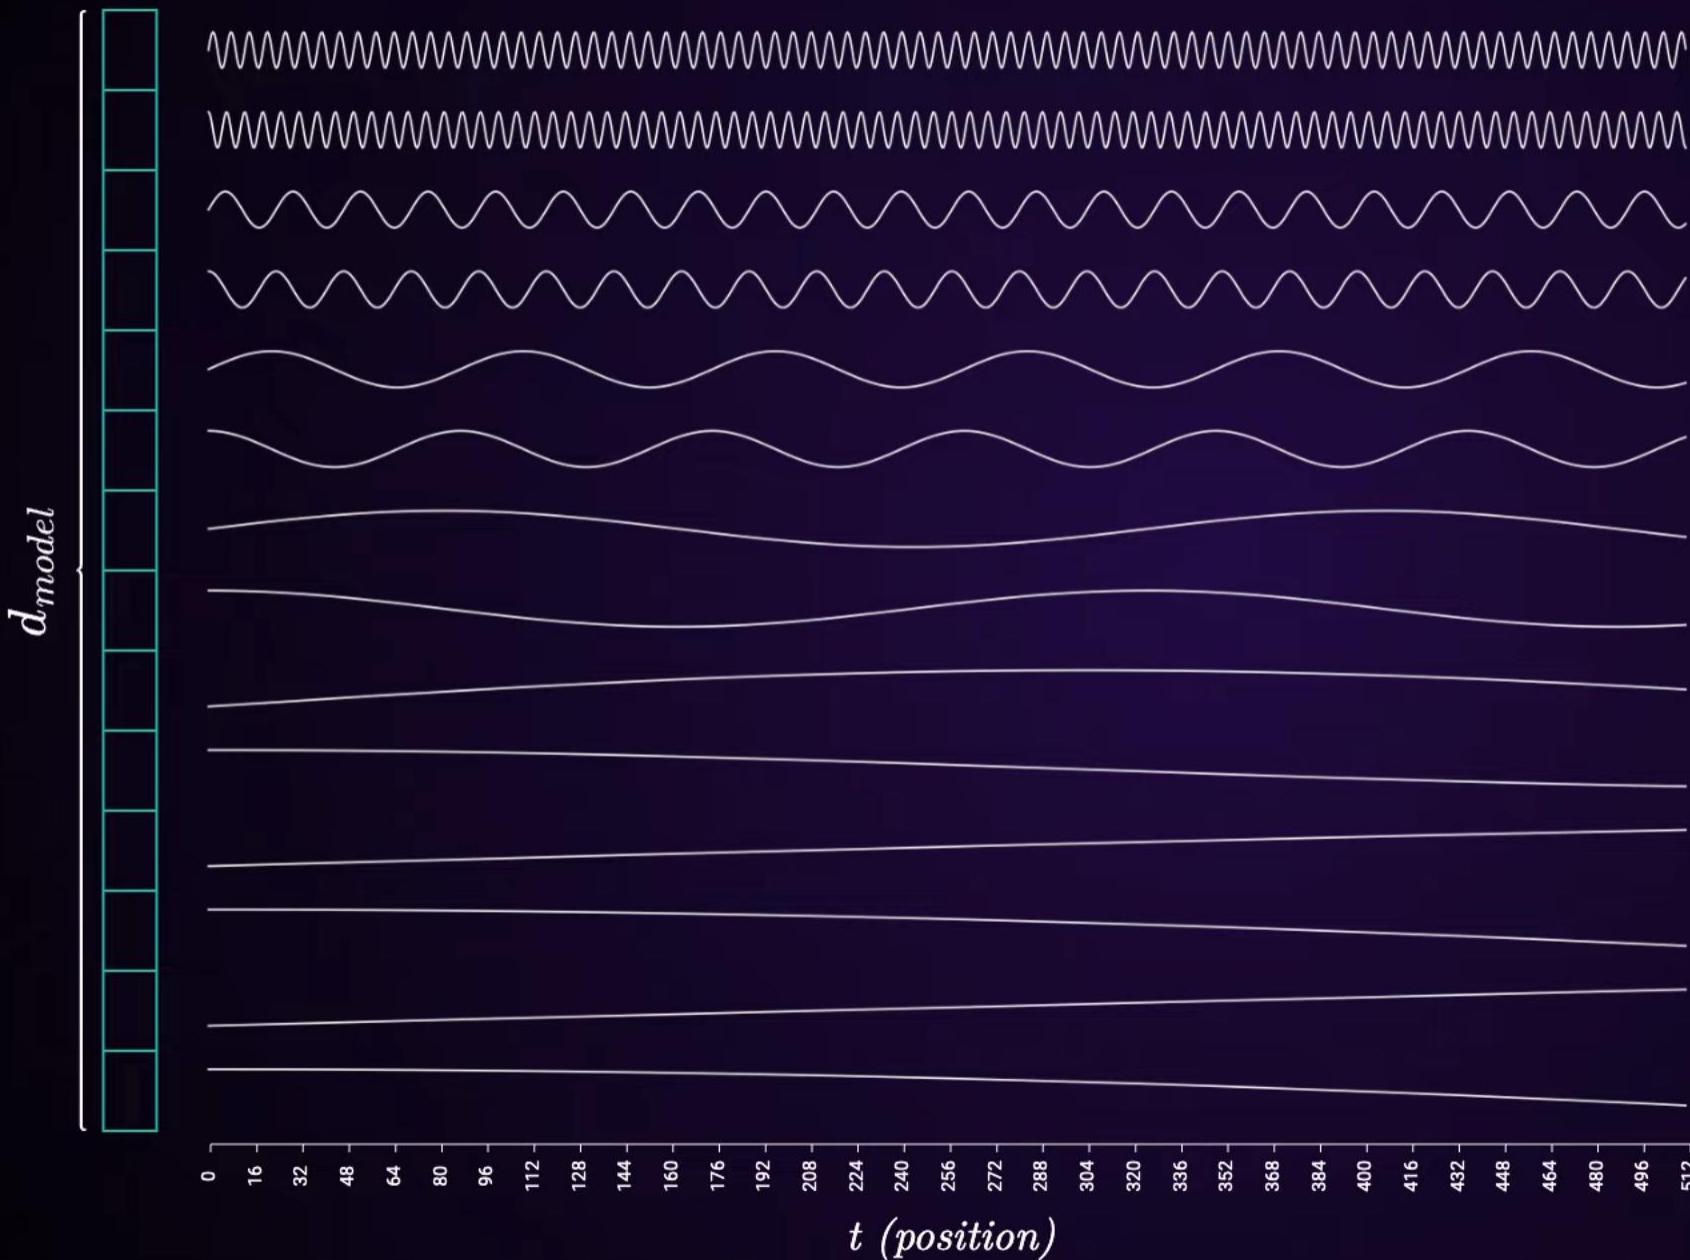
www.youtube.com – To exit full screen, press Esc

$$\vec{v}_t = \begin{bmatrix} \sin(t) \\ \cos(t) \\ \sin(\frac{t}{\eta_1}) \\ \cos(\frac{t}{\eta_1}) \\ \dots \\ \dots \\ \sin(\frac{t}{\eta_{\frac{k}{2}}}) \\ \cos(\frac{t}{\eta_{\frac{k}{2}}}) \end{bmatrix} \quad \vec{v}_{t+\delta} = \begin{bmatrix} \sin(t + \delta) \\ \cos(t + \delta) \\ \sin(\frac{t+\delta}{\eta_1}) \\ \cos(\frac{t+\delta}{\eta_1}) \\ \dots \\ \dots \\ \sin(\frac{t+\delta}{\eta_{\frac{k}{2}}}) \\ \cos(\frac{t+\delta}{\eta_{\frac{k}{2}}}) \end{bmatrix} = \begin{bmatrix} \sin(t)\cos(\delta) + \cos(t)\sin(\delta) \\ \cos(t)\cos(\delta) - \sin(t)\sin(\delta) \\ \sin(\frac{t}{\eta_1})\cos(\frac{\delta}{\eta_1}) + \cos(\frac{t}{\eta_1})\sin(\frac{\delta}{\eta_1}) \\ \cos(\frac{t}{\eta_1})\cos(\frac{\delta}{\eta_1}) - \sin(\frac{t}{\eta_1})\sin(\frac{\delta}{\eta_1}) \\ \dots \\ \dots \\ \sin(\frac{t}{\eta_{\frac{k}{2}}})\cos(\frac{\delta}{\eta_{\frac{k}{2}}}) + \cos(\frac{t}{\eta_{\frac{k}{2}}})\sin(\frac{\delta}{\eta_{\frac{k}{2}}}) \\ \cos(\frac{t}{\eta_{\frac{k}{2}}})\cos(\frac{\delta}{\eta_{\frac{k}{2}}}) - \sin(\frac{t}{\eta_{\frac{k}{2}}})\sin(\frac{\delta}{\eta_{\frac{k}{2}}}) \end{bmatrix} = \begin{bmatrix} \cos(\delta) & \sin(\delta) \\ -\sin(\delta) & \cos(\delta) \\ \cos(\frac{\delta}{\eta_1}) & \sin(\frac{\delta}{\eta_1}) \\ -\sin(\frac{\delta}{\eta_1}) & \cos(\frac{\delta}{\eta_1}) \\ \dots \\ \cos(\frac{\delta}{\eta_{\frac{k}{2}}}) & \sin(\frac{\delta}{\eta_{\frac{k}{2}}}) \\ -\sin(\frac{\delta}{\eta_{\frac{k}{2}}}) & \cos(\frac{\delta}{\eta_{\frac{k}{2}}}) \end{bmatrix} \times \begin{bmatrix} \sin(t) \\ \cos(t) \\ \sin(\frac{t}{\eta_1}) \\ \cos(\frac{t}{\eta_1}) \\ \dots \\ \dots \\ \sin(\frac{t}{\eta_{\frac{k}{2}}}) \\ \cos(\frac{t}{\eta_{\frac{k}{2}}}) \end{bmatrix}$$

$\sin(\alpha + \beta) = \sin(\alpha)\cos(\beta) + \cos(\alpha)\sin(\beta)$
 $\cos(\alpha + \beta) = \cos(\alpha)\cos(\beta) - \sin(\alpha)\sin(\beta)$

- For any fixed offset k , PE($t+k$) can be represented as a linear function of PE(t).

<https://www.youtube.com/watch?v=T3OT8kqoqjc>



$$k \text{ is even: } \sin\left(\frac{t}{N^{\frac{k}{d_{model}}}}\right)$$

$$k \text{ is odd: } \cos\left(\frac{t}{N^{\frac{k}{d_{model}}}}\right)$$

$$L \text{ (context/sequence length)} = 512$$

$$t \text{ (position)} = \{1, \dots, L\}$$

$$N = 10000$$

$$d_{model} = 14$$

$$k = \{1, 2, \dots, d_{model}\}$$

Not Really that Useful...



```
54
55     class BertEmbeddings(nn.Module):
56         """Construct the embeddings from word, position and token_type embeddings."""
57
58     def __init__(self, config):
59         super().__init__()
60         self.word_embeddings = nn.Embedding(config.vocab_size, config.hidden_size, padding_idx=config.
61         self.position_embeddings = nn.Embedding(config.max_position_embeddings, config.hidden_size)
62         self.token_type_embeddings = nn.Embedding(config.type_vocab_size, config.hidden_size)
63
64         self.LayerNorm = nn.LayerNorm(config.hidden_size, eps=config.layer_norm_eps)
65         self.dropout = nn.Dropout(config.hidden_dropout_prob)
66         # position_ids (1, len position emb) is contiguous in memory and exported when serialized
67         self.register_buffer(
68             "position_ids", torch.arange(config.max_position_embeddings).expand((1, -1)), persistent=False
69         )
70         self.register_buffer(
71             "token_type_ids", torch.zeros(self.position_ids.size(), dtype=torch.long), persistent=False
72         )
73
```

BERT

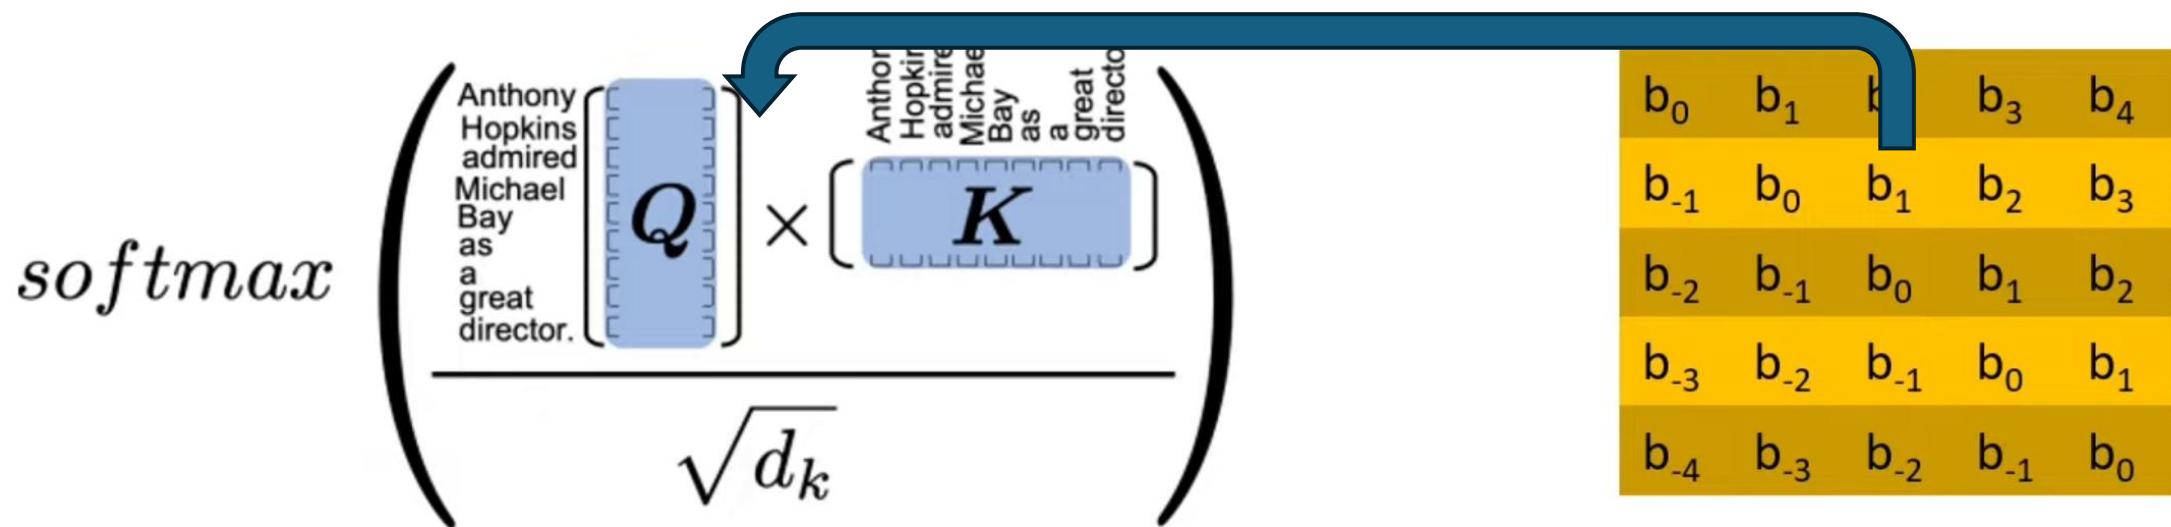
```
544
545     @auto_docstring
546     class GPT2Model(GPT2PreTrainedModel):
547         def __init__(self, config):
548             super().__init__(config)
549
550             self.embed_dim = config.hidden_size
551
552             self.wte = nn.Embedding(config.vocab_size, self.embed_dim)
553             self.wpe = nn.Embedding(config.max_position_embeddings, self.embed_dim)
554
555             self.drop = nn.Dropout(config.embd_pdrop)
556             self.h = nn.ModuleList([GPT2Block(config, layer_idx=i) for i in range(config.num_hidden_layers)])
557             self.ln_f = nn.LayerNorm(self.embed_dim, eps=config.layer_norm_epsilon)
558
559             self.gradient_checkpointing = False
560             self._attnImplementation = config._attnImplementation
561
562             # Initialize weights and apply final processing
563             self.post_init()
```

GPT-2

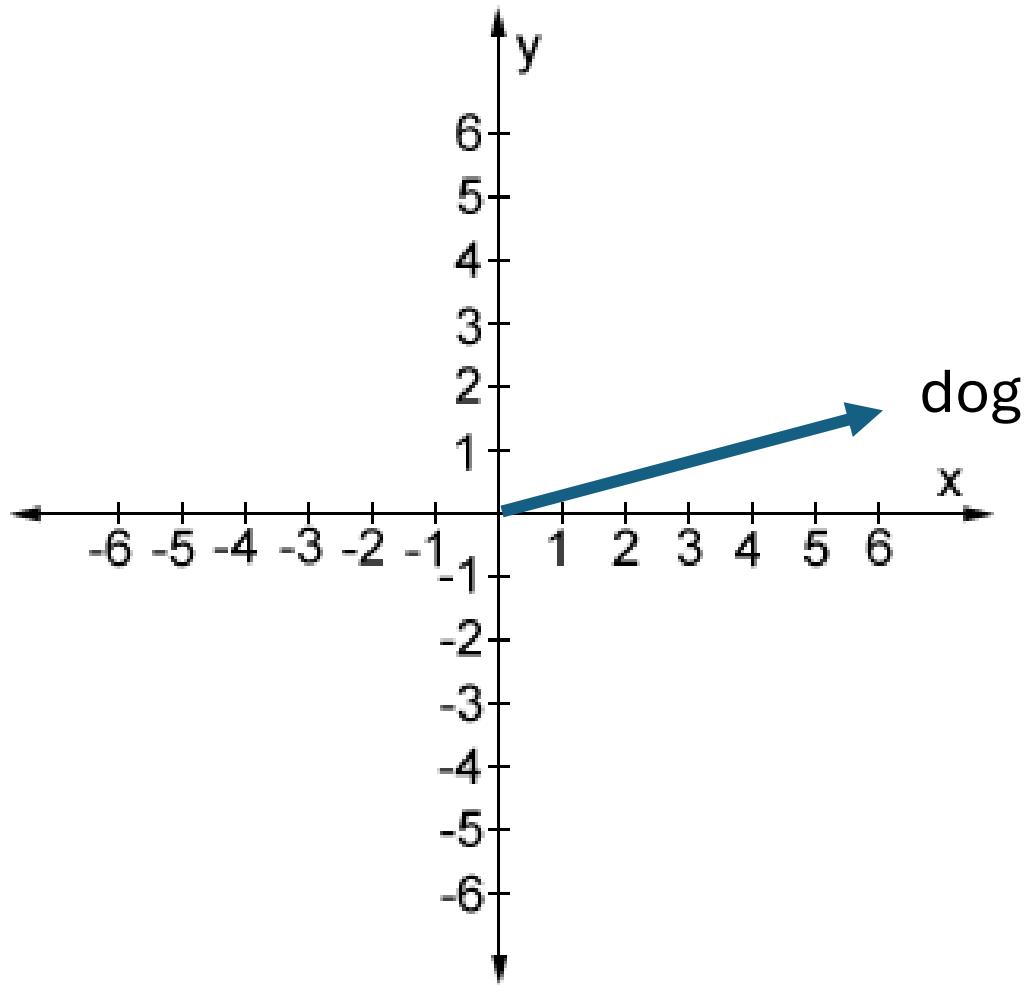
Just treat positions 0,1,2,3, ..., N as “tokens.”

Relative Position Embedding: T5

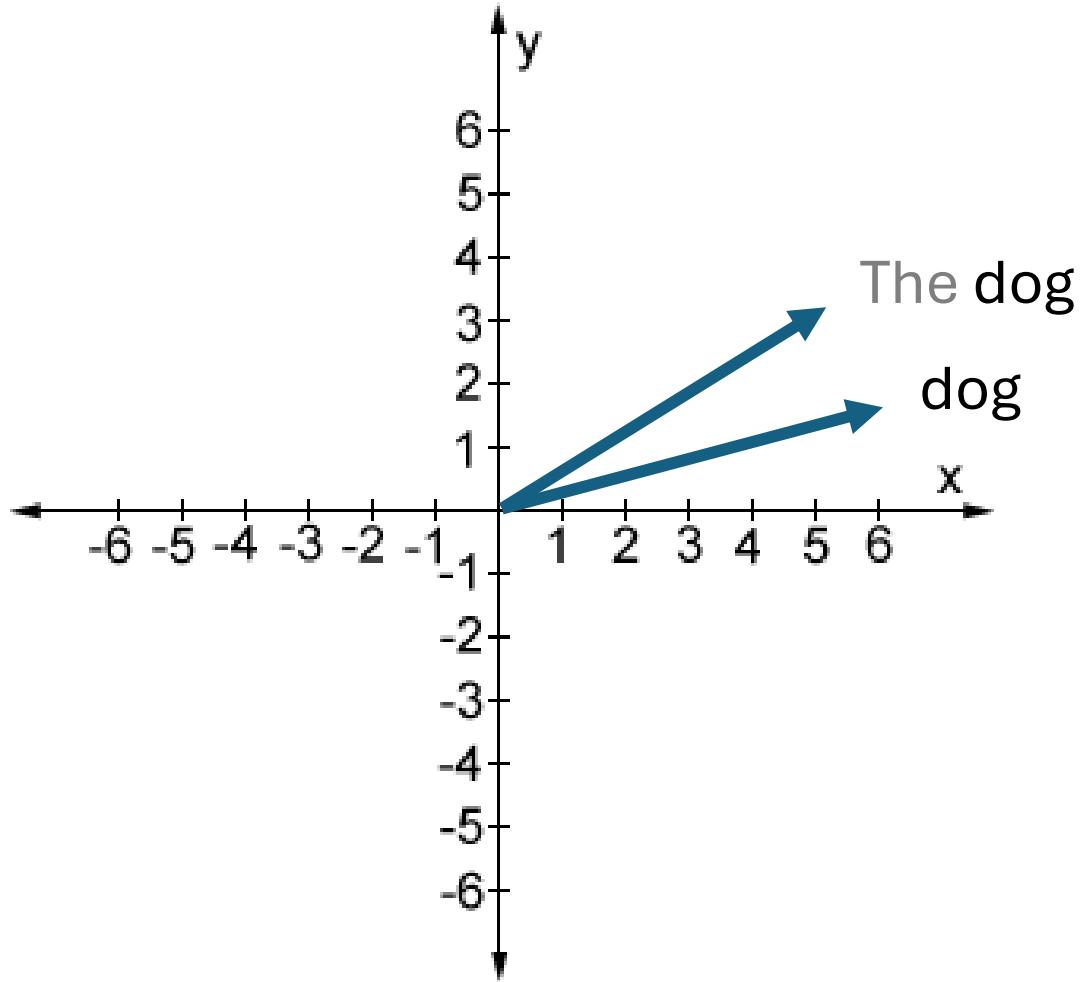
- Absolute position embedding:
 - One position, one embedding. $E = E_{\text{token}} + E_{\text{pos}}$.
- Relative position embedding:



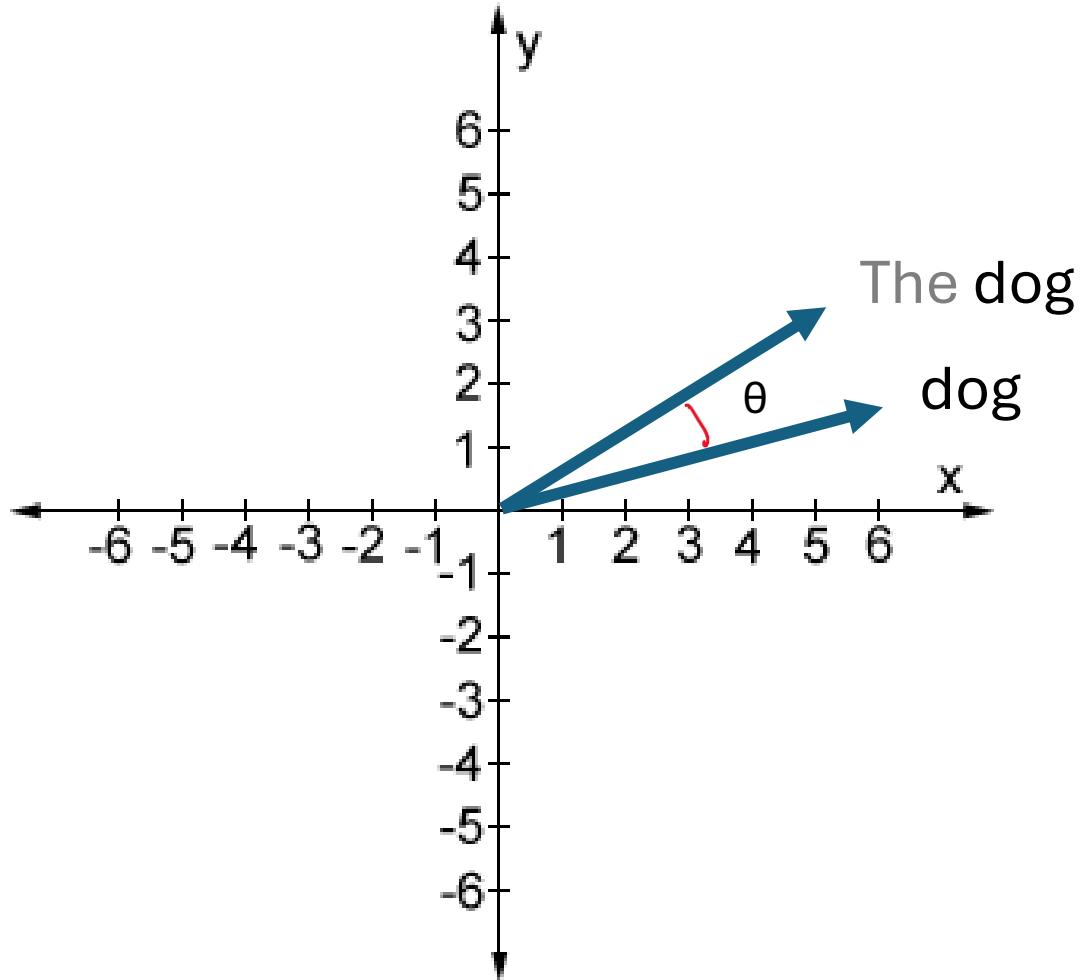
RoPE: Rotary Position Embedding



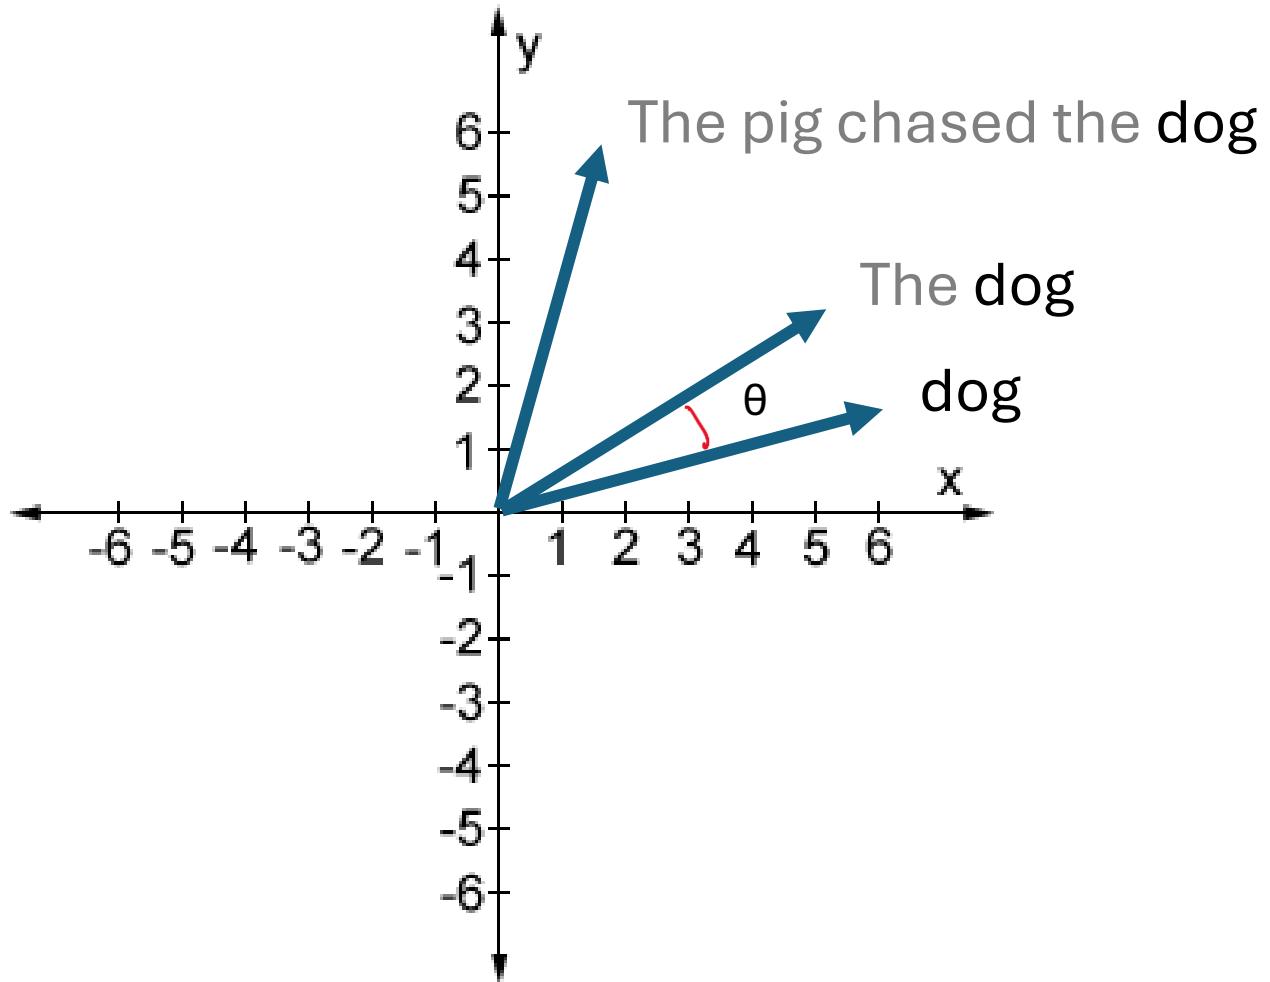
RoPE: Rotary Position Embedding



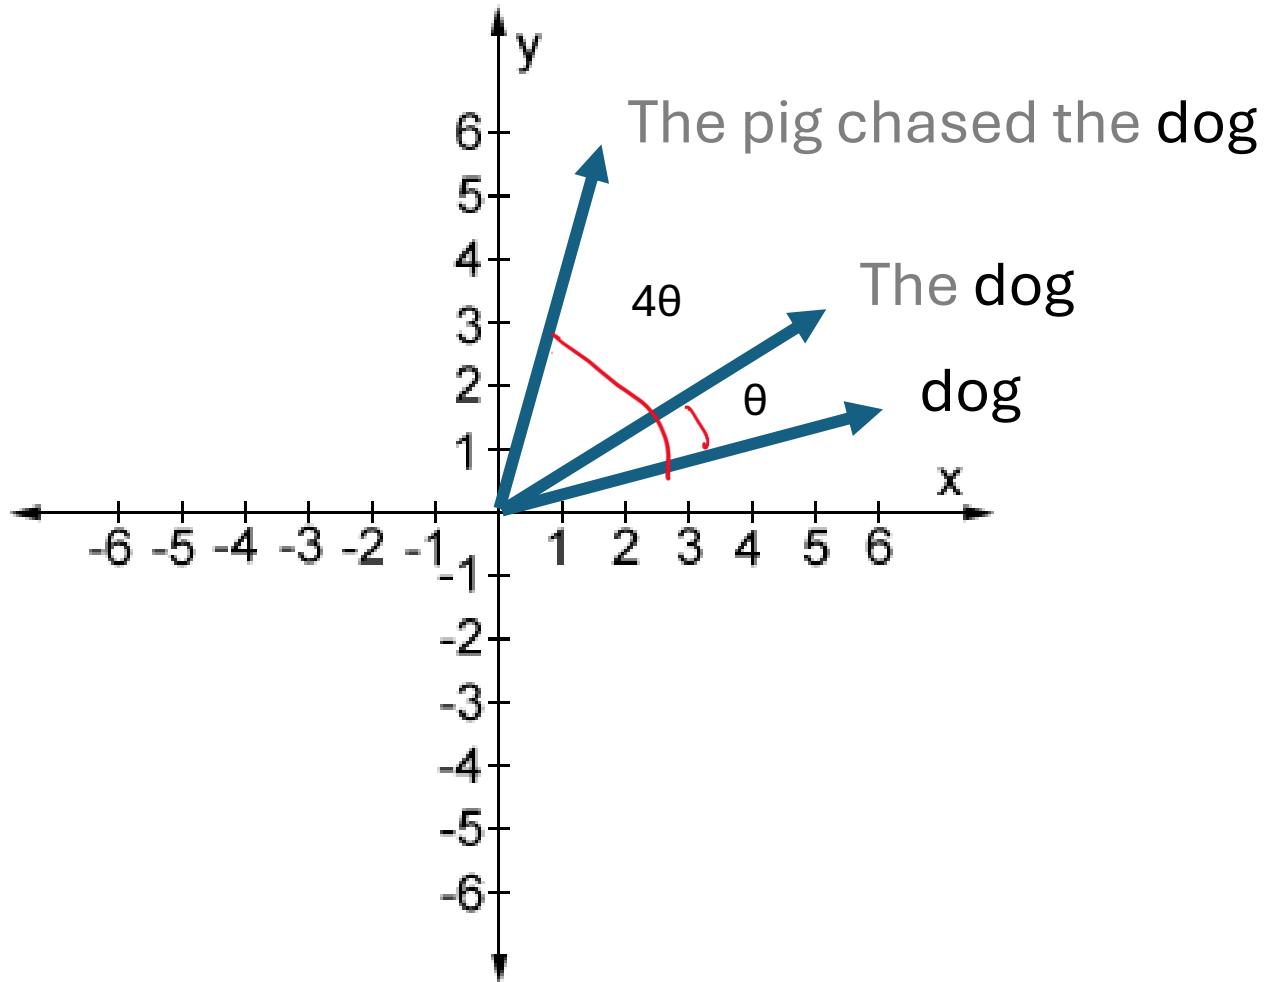
RoPE: Rotary Position Embedding



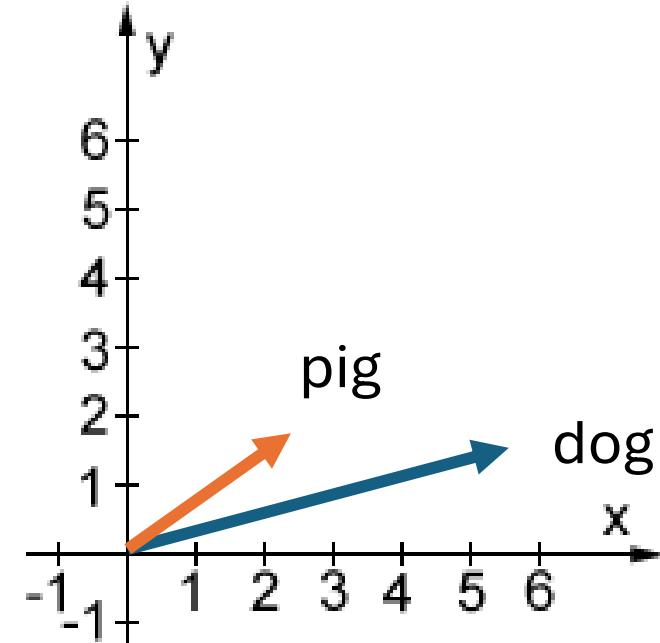
RoPE: Rotary Position Embedding



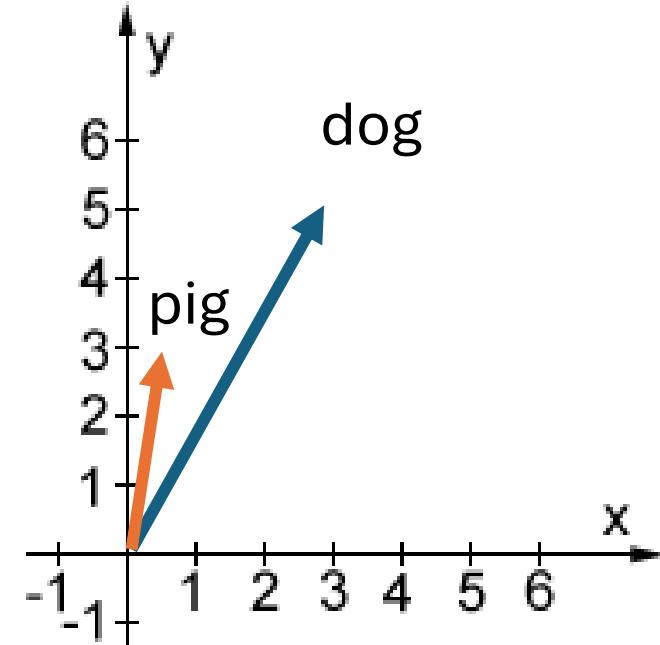
RoPE: Rotary Position Embedding



The pig chased the dog



Once upon a time,
the pig chased the dog



Implementation: 2D Example

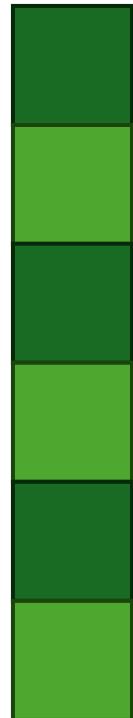
$$f_{\{q,k\}}(\mathbf{x}_m, m) = \begin{pmatrix} \cos m\theta & -\sin m\theta \\ \sin m\theta & \cos m\theta \end{pmatrix} \begin{pmatrix} W_{\{q,k\}}^{(11)} \\ W_{\{q,k\}}^{(21)} \\ W_{\{q,k\}}^{(22)} \end{pmatrix} \begin{pmatrix} \mathbf{x}_m^{(1)} \\ \mathbf{x}_m^{(2)} \end{pmatrix}$$


Rotate a vector by $m\theta$

Implementation: General Case

$$f_{\{q,k\}}(\mathbf{x}_m, m) = \mathbf{R}_{\Theta,m}^d \mathbf{W}_{\{q,k\}} \mathbf{x}_m$$

$$\mathbf{R}_{\Theta,m}^d = \begin{pmatrix} \cos m\theta_1 & -\sin m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ \sin m\theta_1 & \cos m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 0 & \cos m\theta_2 & -\sin m\theta_2 & \cdots & 0 & 0 \\ 0 & 0 & \sin m\theta_2 & \cos m\theta_2 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & \cos m\theta_{d/2} & -\sin m\theta_{d/2} \\ 0 & 0 & 0 & 0 & \cdots & \sin m\theta_{d/2} & \cos m\theta_{d/2} \end{pmatrix}$$



Implementation: Actually...

```
def rotate_half(x):
    """Rotates half the hidden dims of the input."""
    x1 = x[..., : x.shape[-1] // 2]
    x2 = x[..., x.shape[-1] // 2 :]
    return torch.cat((-x2, x1), dim=-1)

def apply_rotary_pos_emb(q, k, cos, sin, position_ids=None, unsqueeze_dim=1):
    """Applies Rotary Position Embedding to the query and key tensors."""
    cos = cos.unsqueeze(unsqueeze_dim)
    sin = sin.unsqueeze(unsqueeze_dim)
    q_embed = (q * cos) + (rotate_half(q) * sin)
    k_embed = (k * cos) + (rotate_half(k) * sin)
    return q_embed, k_embed
```

RoPE Summary

- Su et al., (2021). RoFormer: Enhanced Transformer with Rotary Position Embedding.
- Good performance:
- Combines the strength of absolute and relative position encoding.
- The most widely used type of position encoding in contemporary LMs.
 - Basically all major models right now: LLaMA, Mistral, Qwen, GPT-NeoX, Falcon, Deepseek, Gemma...
 - ChatGPT (GPT-4/5), Grok, Claude, ... probably also use RoPE.

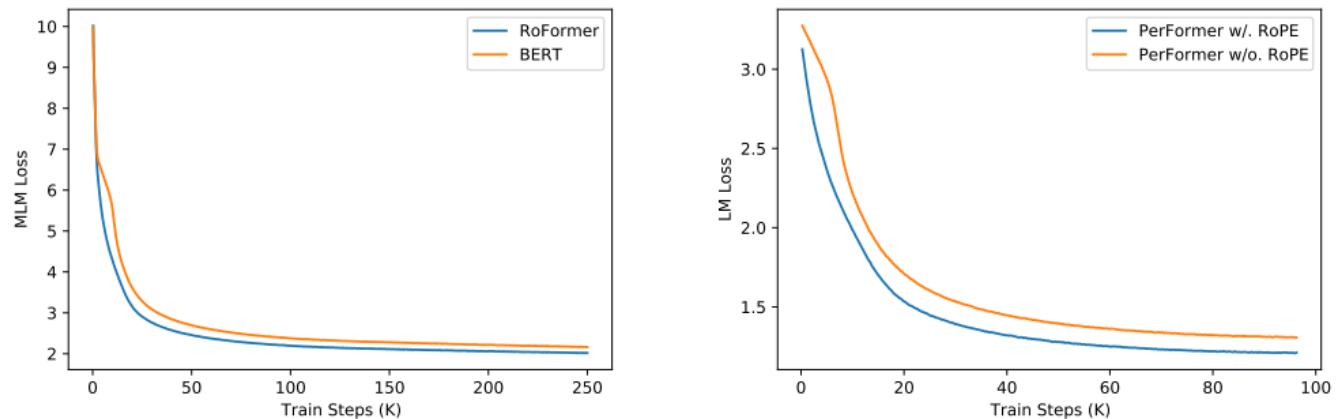
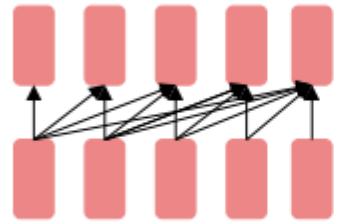


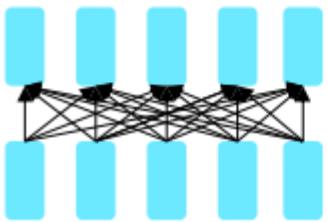
Figure 3: Evaluation of RoPE in language modeling pre-training. **Left:** training loss for BERT and RoFormer. **Right:** training loss for PerFormer with and without RoPE.

Three types of architectures



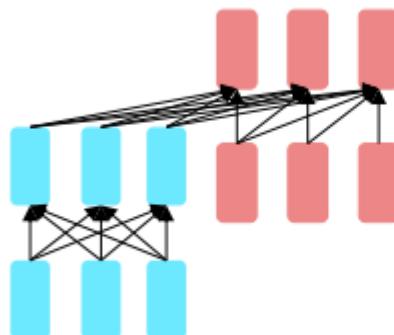
Decoders

- Next word prediction.
- Easy to train. Abundant amount of data.
- Nice to generate from; can't condition on future words.



Encoders

- Gets bidirectional context – can condition on future!
- Good word embeddings.
- MLM, BERT.

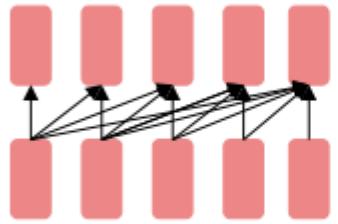


Encoder-Decoders

- Good parts of decoders and encoders?
- What's the best way to pretrain them?

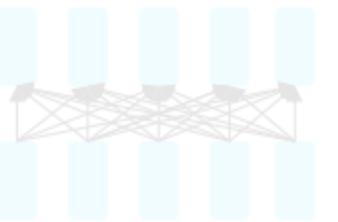
We will come to more details...

Three types of architectures



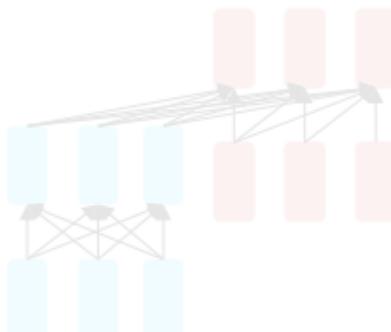
Decoders

- Next word prediction.
- Easy to train. Abundant amount of data.
- Nice to generate from; can't condition on future words.



Encoders

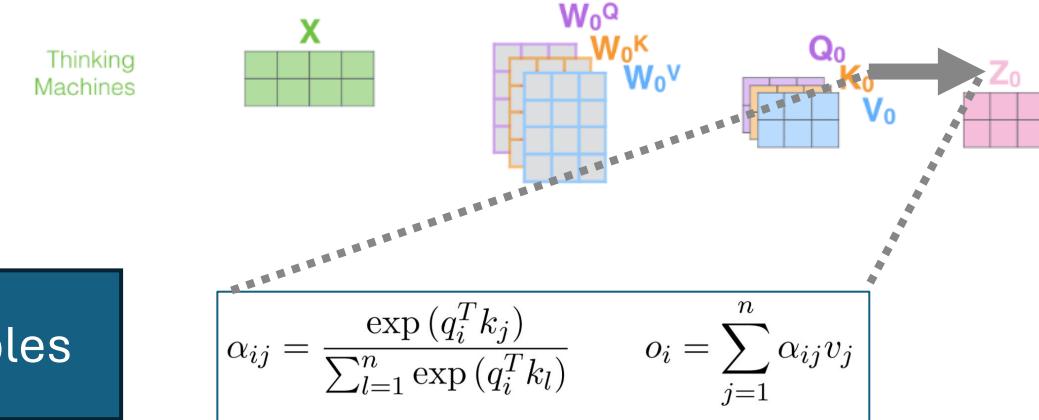
- Gets bidirectional context – can condition on future!
- Good word embeddings.
- MLM, BERT.



Encoder-Decoders

- Good parts of decoders and encoders?
- What's the best way to pretrain them?

Self Attention



I like eating apples

What is $o(\text{like})$? i.e., the attention score from like to the sentence.

query

Q4
“apples”

key

K1 “I”	K2 “like”	K3 “eating”	K4 “apples”
-----------	--------------	----------------	----------------

attention
weights*

14.2	18.1	10.3	7.9
------	------	------	-----

attention
scores* (o)

0.3	0.6	0.07	0.03
-----	-----	------	------

*: made-up numbers, not real.

$$X \times W^Q = Q$$

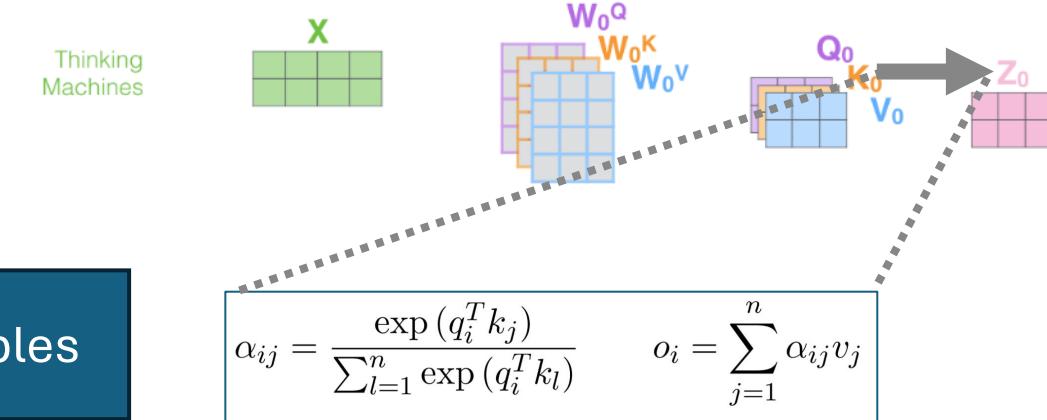
$$X \times W^K = K$$

$$Q \times K^T$$

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}} \right)$$

Self Attention

I like eating apples



What is $z(\text{like})$?

attention scores* (o) 0.3 0.6 0.07 0.03

values

V1 “I”	V2 “like”	V3 “eating”	V4 “apples”
-----------	--------------	----------------	----------------

Weighted sum

Z4
“like”

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}} \right)$$

$$X \times W^V = V$$

*: made-up numbers, not real.

GPT-2

Language Models are Unsupervised Multitask Learners

	LAMBADA (PPL)	LAMBADA (ACC)	CBT-CN (ACC)	CBT-NE (ACC)	WikiText2 (PPL)	PTB (PPL)	enwik8 (BPB)	text8 (BPC)	WikiText103 (PPL)	1BW (PPL)
SOTA	99.8	59.23	85.7	82.3	39.14	46.54	0.99	1.08	18.3	21.8
117M	35.13	45.99	87.65	83.4	29.41	65.85	1.16	1.17	37.50	75.20
345M	15.60	55.48	92.35	87.1	22.76	47.33	1.01	1.06	26.37	55.72
762M	10.87	60.12	93.45	88.0	19.93	40.31	0.97	1.02	22.05	44.575
1542M	8.63	63.24	93.30	89.05	18.34	35.76	0.93	0.98	17.48	42.16

Table 3. Zero-shot results on many datasets. No training or fine-tuning was performed for any of these results. PTB and WikiText-2 results are from (Gong et al., 2018). CBT results are from (Bajgar et al., 2016). LAMBADA accuracy result is from (Hoang et al., 2018) and LAMBADA perplexity result is from (Grave et al., 2016). Other results are from (Dai et al., 2019).

Perplexity

$$\text{PPL}(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$

Normalized by the number of words.



Inverse probability of a corpus, according to a LM.

Perplexity... Boring!

OpenAI's GPT-2 has been promoted as "an AI" that exemplifies an emergent understanding of language after mere unsupervised training on about 40GB of webpage text. It sounds really convincing in interviews:

- Q: *Which technologies are worth watching in 2020?*

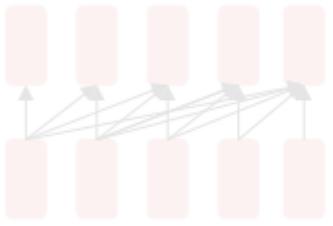
A: I would say it is hard to narrow down the list. The world is full of disruptive technologies with real and potentially huge global impacts. The most important is artificial intelligence, which is becoming exponentially more powerful. There is also the development of self-driving cars. There is a lot that we can do with artificial intelligence to improve the world....

- Q: *Are you worried that ai [sic] technology can be misused?*

A: Yes, of course. But this is a global problem and we want to tackle it with global solutions.....

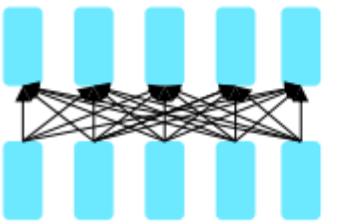
--- "AI can do that", *The World in 2020 – The Economist*

Three types of architectures



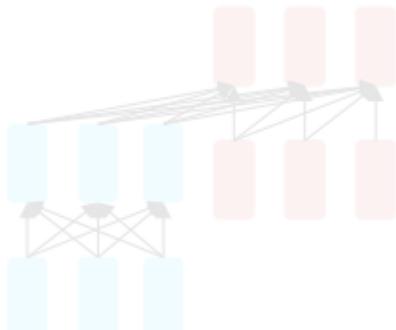
Decoders

- Next word prediction.
- Easy to train. Abundant amount of data.
- Nice to generate from; can't condition on future words.



Encoders

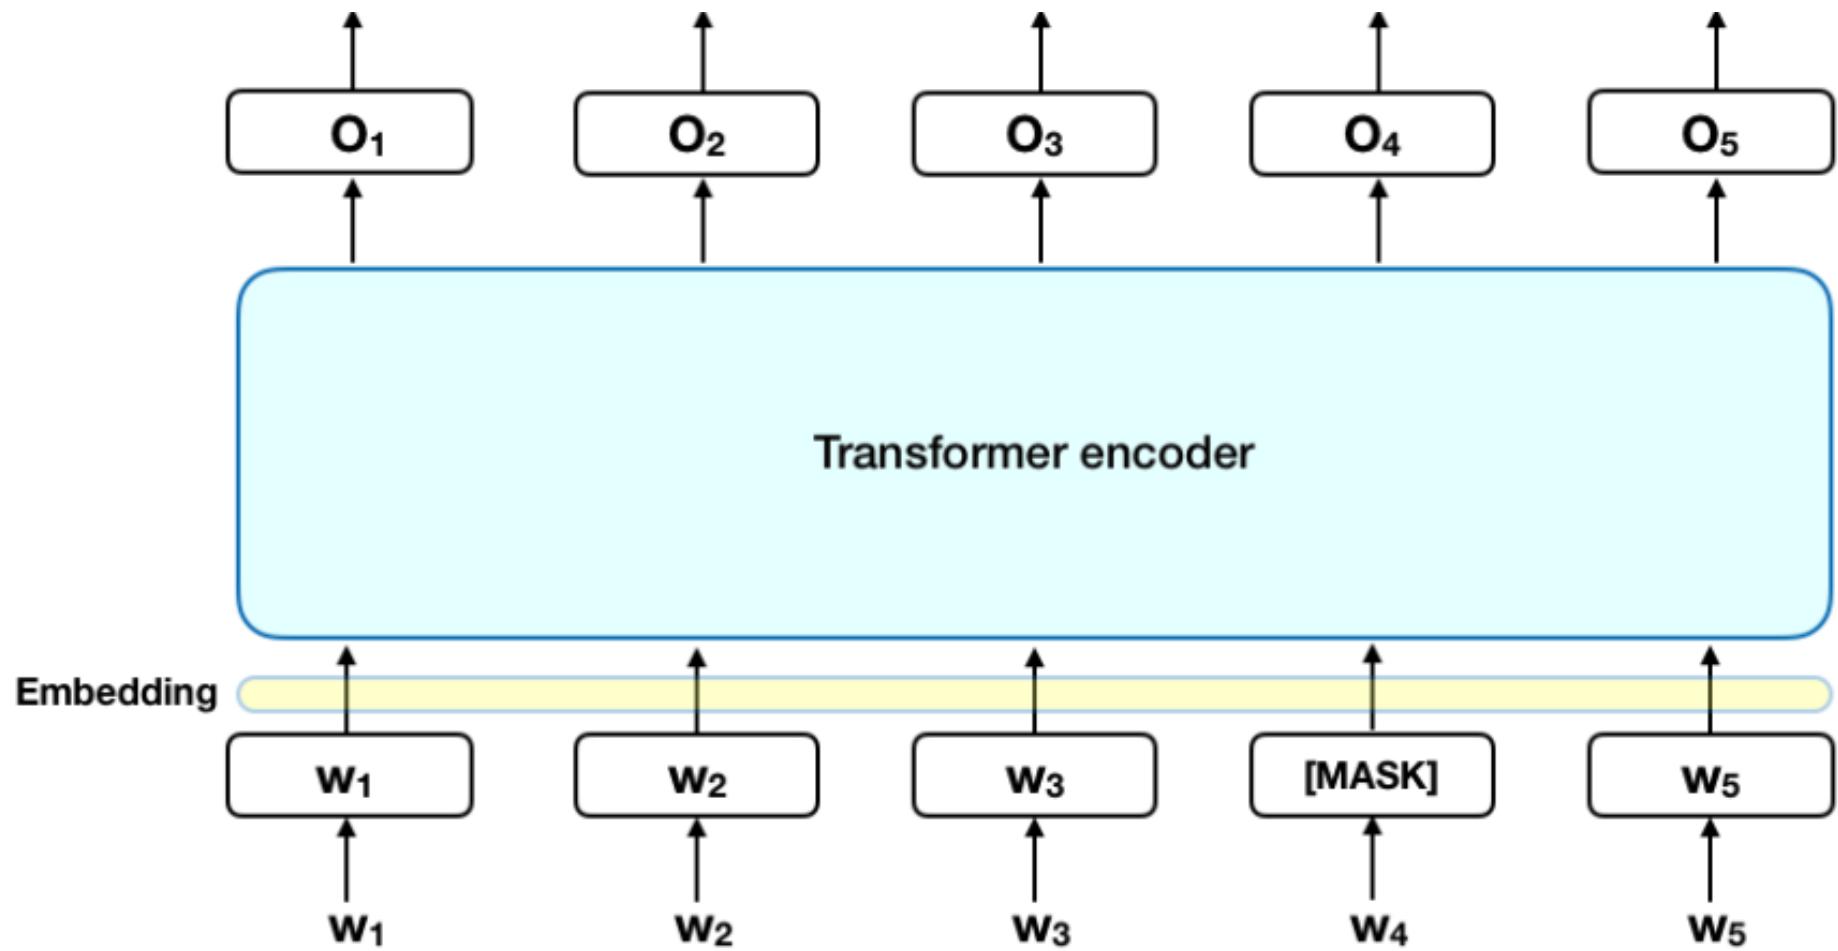
- Gets bidirectional context – can condition on future!
- Good word embeddings.
- MLM, BERT.



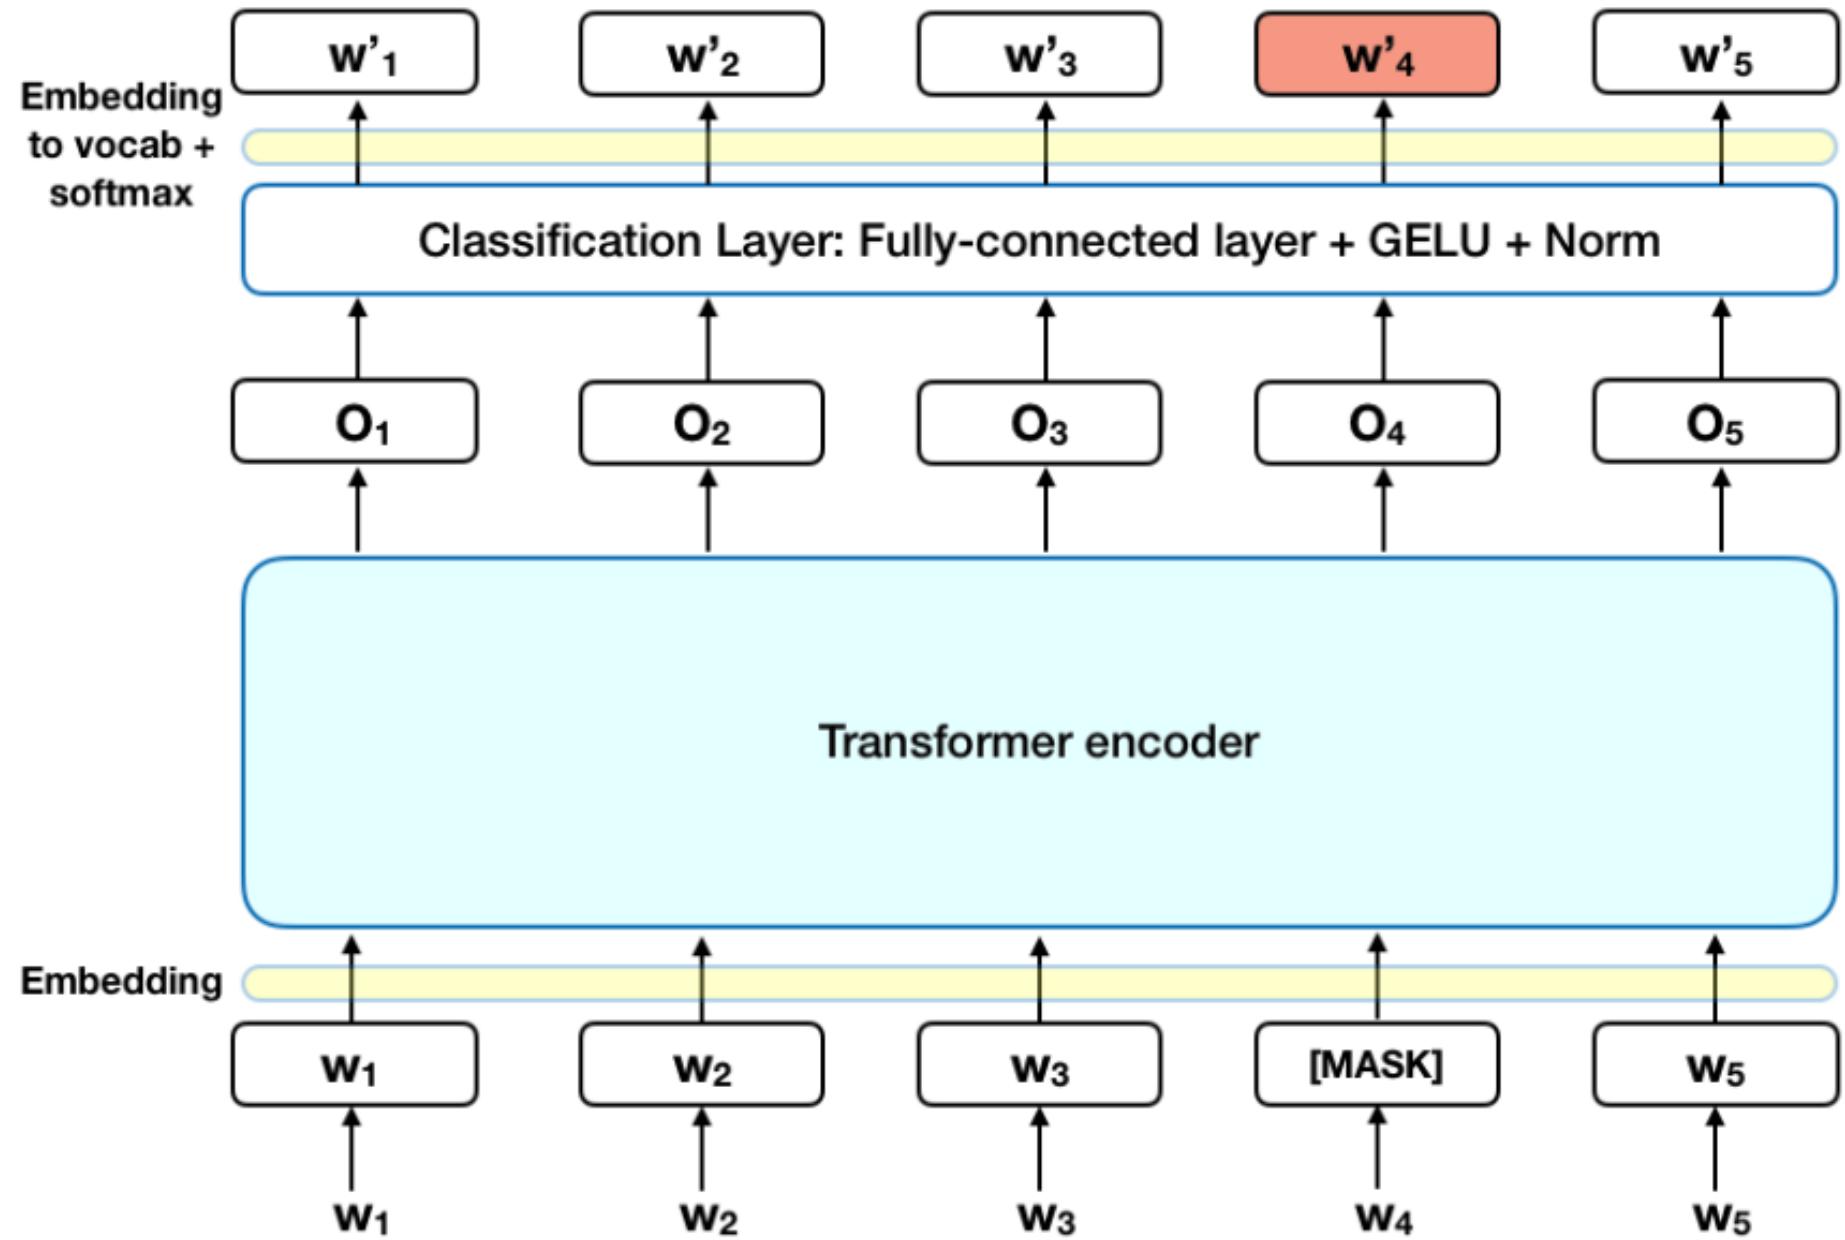
Encoder-Decoders

- Good parts of decoders and encoders?
- What's the best way to pretrain them?

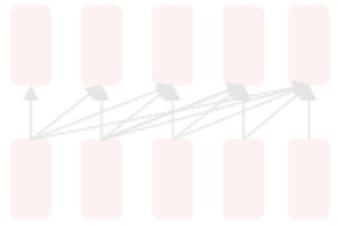
BERT



BERT

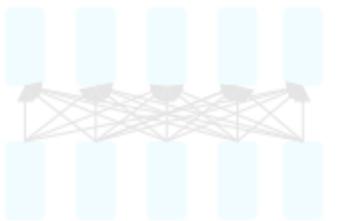


Three types of architectures



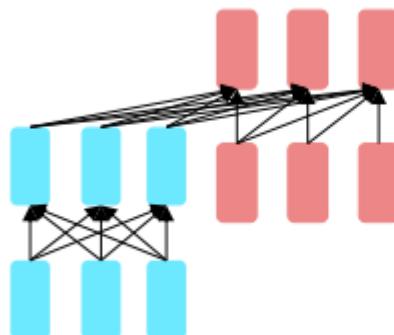
Decoders

- Next word prediction.
- Easy to train. Abundant amount of data.
- Nice to generate from; can't condition on future words.



Encoders

- Gets bidirectional context – can condition on future!
- Good word embeddings.
- MLM, BERT.



Encoder-Decoders

- Good parts of decoders and encoders?
- What's the best way to pretrain them?

Pretraining encoder-decoders: What pretraining objective to use?



- What Raffel et al. (2018) found to work best was span corruption. Their model: T5.
- Replace different-length spans from the input with unique placeholders; decode out the spans that were removed!

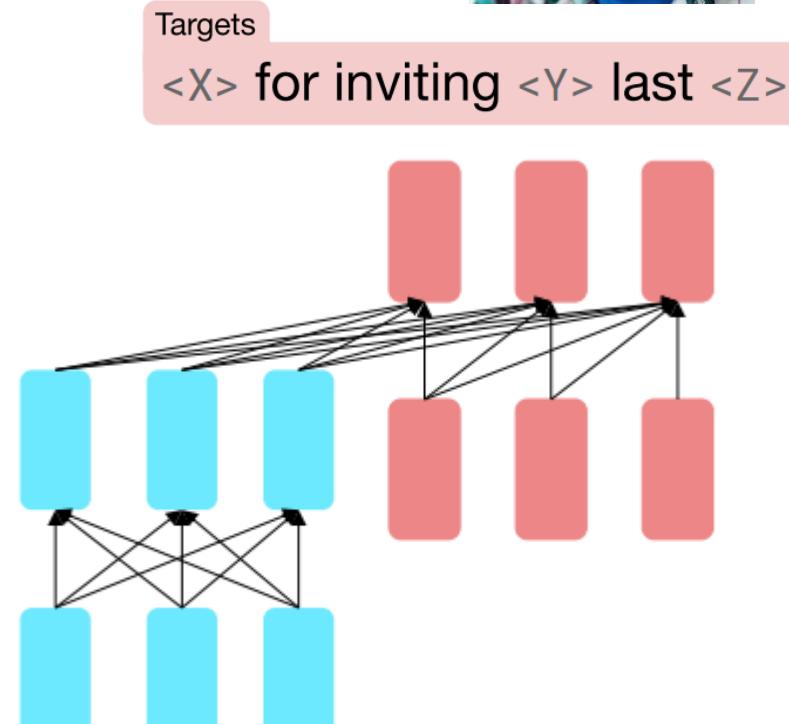
Original text

Thank you ~~for~~ ~~inviting~~ me to your party ~~last~~ week.

- This is implemented in text preprocessing: it's still an objective that looks like language modeling at the decoder side.

Inputs

Thank you <X> me to your party <Y> week.

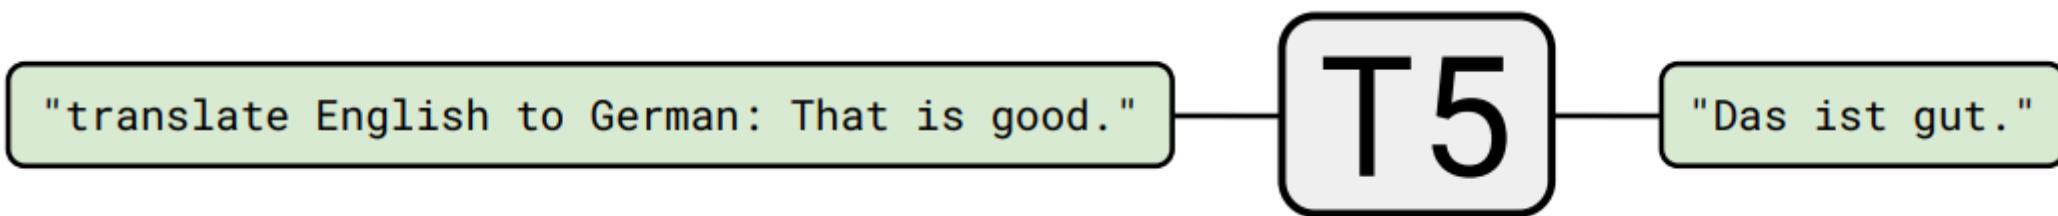


Pretraining encoder-decoders: What pretraining objective to use?

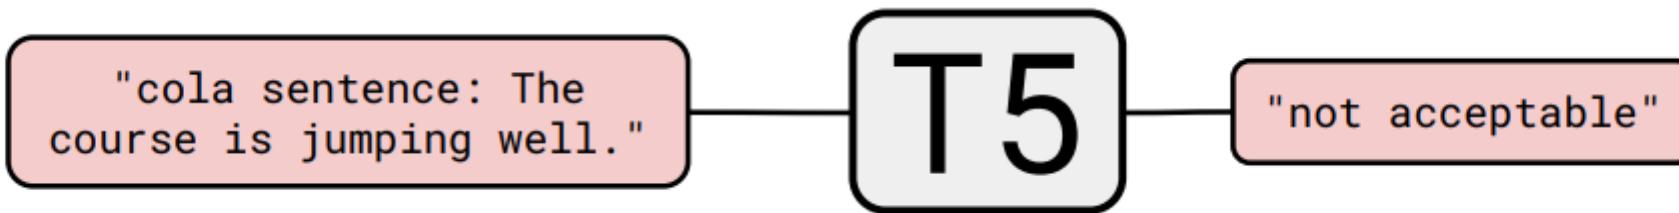
- Raffel et al., (2018) found encoder-decoders to work better than decoders for their tasks, and span corruption (denoising) to work better than language modeling.

Architecture	Objective	Params	Cost	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Encoder-decoder	Denoising	$2P$	M	83.28	19.24	80.88	71.36	26.98	39.82	27.65
	Denoising	P	M	82.81	18.78	80.63	70.73	26.72	39.03	27.46
	Denoising	P	$M/2$	80.88	18.97	77.59	68.42	26.38	38.40	26.95
	Denoising	P	M	74.70	17.93	61.14	55.02	25.09	35.28	25.86
	Denoising	P	M	81.82	18.61	78.94	68.11	26.43	37.98	27.39
Encoder-decoder	LM	$2P$	M	79.56	18.59	76.02	64.29	26.27	39.17	26.86
	LM	P	M	79.60	18.13	76.35	63.50	26.62	39.17	27.05
	LM	P	$M/2$	78.67	18.26	75.32	64.06	26.13	38.42	26.89
	LM	P	M	73.78	17.54	53.81	56.51	25.23	34.31	25.38
	LM	P	M	79.68	17.84	76.87	64.86	26.28	37.51	26.76

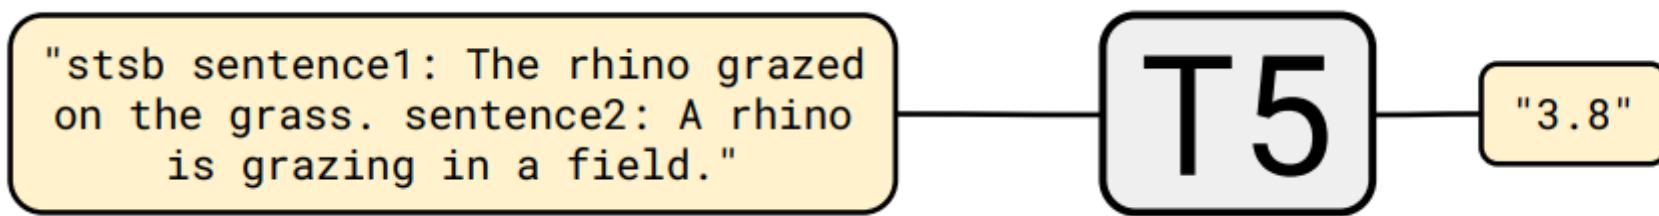
One surprising finding



One surprising finding



One surprising finding



One surprising finding

"translate English to German: That is good."

"cola sentence: The course is jumping well."

"stsbs sentence1: The rhino grazed on the grass. sentence2: A rhino is grazing in a field."

"summarize: state authorities dispatched emergency crews tuesday to survey the damage after an onslaught of severe weather in mississippi..."

A fascinating property of T5:

- It can be finetuned to answer a wide range of questions, retrieving knowledge from its parameters.
- With natural language!

T5

"Das ist gut."

"not acceptable"

"3.8"

"six people hospitalized after a storm in attala county."

One surprising finding

Text-to-Text Transfer Transformer

"translate English to German: That is good."

"cola sentence: The course is jumping well."

"stsbs sentence1: The rhino grazed on the grass. sentence2: A rhino is grazing in a field."

"summarize: state authorities dispatched emergency crews tuesday to survey the damage after an onslaught of severe weather in mississippi..."

T5

"Das ist gut."

"not acceptable"

"3.8"

"six people hospitalized after a storm in attala county."

Input

Embedding

Queries

Keys

Values

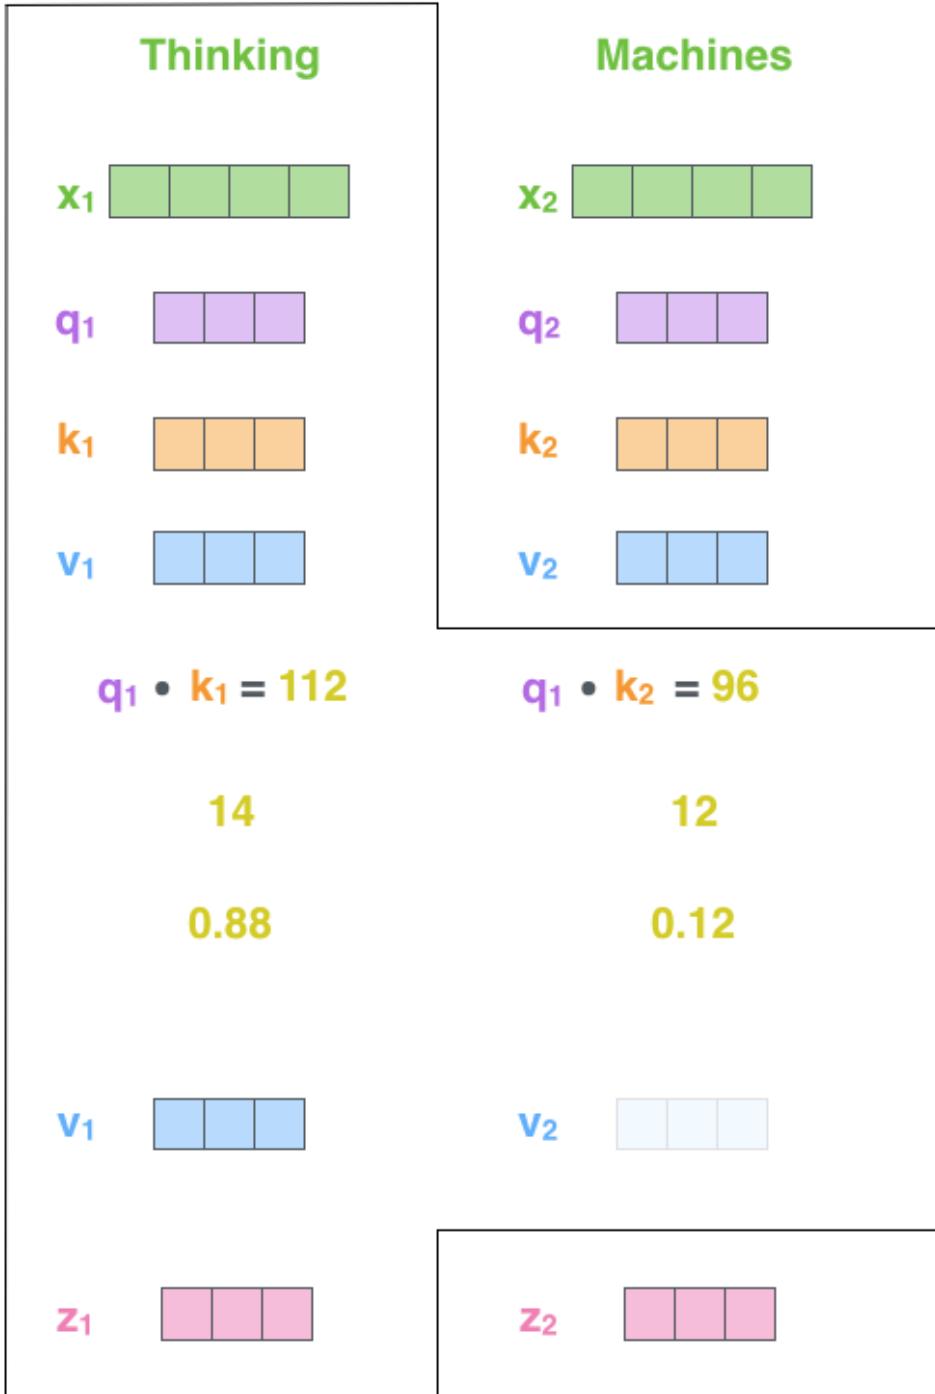
Score

Divide by 8 ($\sqrt{d_k}$)

Softmax

Softmax
X
Value

Sum



Review

- Embeddings:
 - Token + Sentence + Position
- Multi-Head Attention
- Feed forward module (MLP)
- Layers